

# **INNOVATION LAB**

**(TITLE: GROCERY OBJECT DETECTION: IMPROVING MODEL  
PERFORMANCE THROUGH DATA AUGMENTATION)**



**NAME: HIMANI YADAV  
ROLL NO: 2201AI16  
UNDER - DR.JIMSON MATHEW**



- **OBJECTIVE:**

- **BUILD A GROCERY OBJECT DETECTION MODEL.**
- **ENHANCE PERFORMANCE USING DATA AUGMENTATION TECHNIQUES.**

- **CHALLENGES IDENTIFIED:**

- **POOR MODEL PERFORMANCE ON INITIAL TRAINING.**

# **OBJECT DETECTION**

## **USING YOLOV7**

- DATASET ANNOTATION**
- MODEL TRAINING**
- MODEL TESTING**

# INTRODUCTION TO OBJECT DETECTION

- OBJECT DETECTION IS THE TASK OF IDENTIFYING OBJECTS WITHIN AN IMAGE OR VIDEO.
- IT OUTPUTS BOTH CLASS LABELS AND BOUNDING BOXES FOR EACH OBJECT DETECTED.
- WIDELY USED IN APPLICATIONS LIKE AUTONOMOUS DRIVING, SECURITY, MEDICAL IMAGING, ETC.



# INTRODUCTION TO YOLOV7

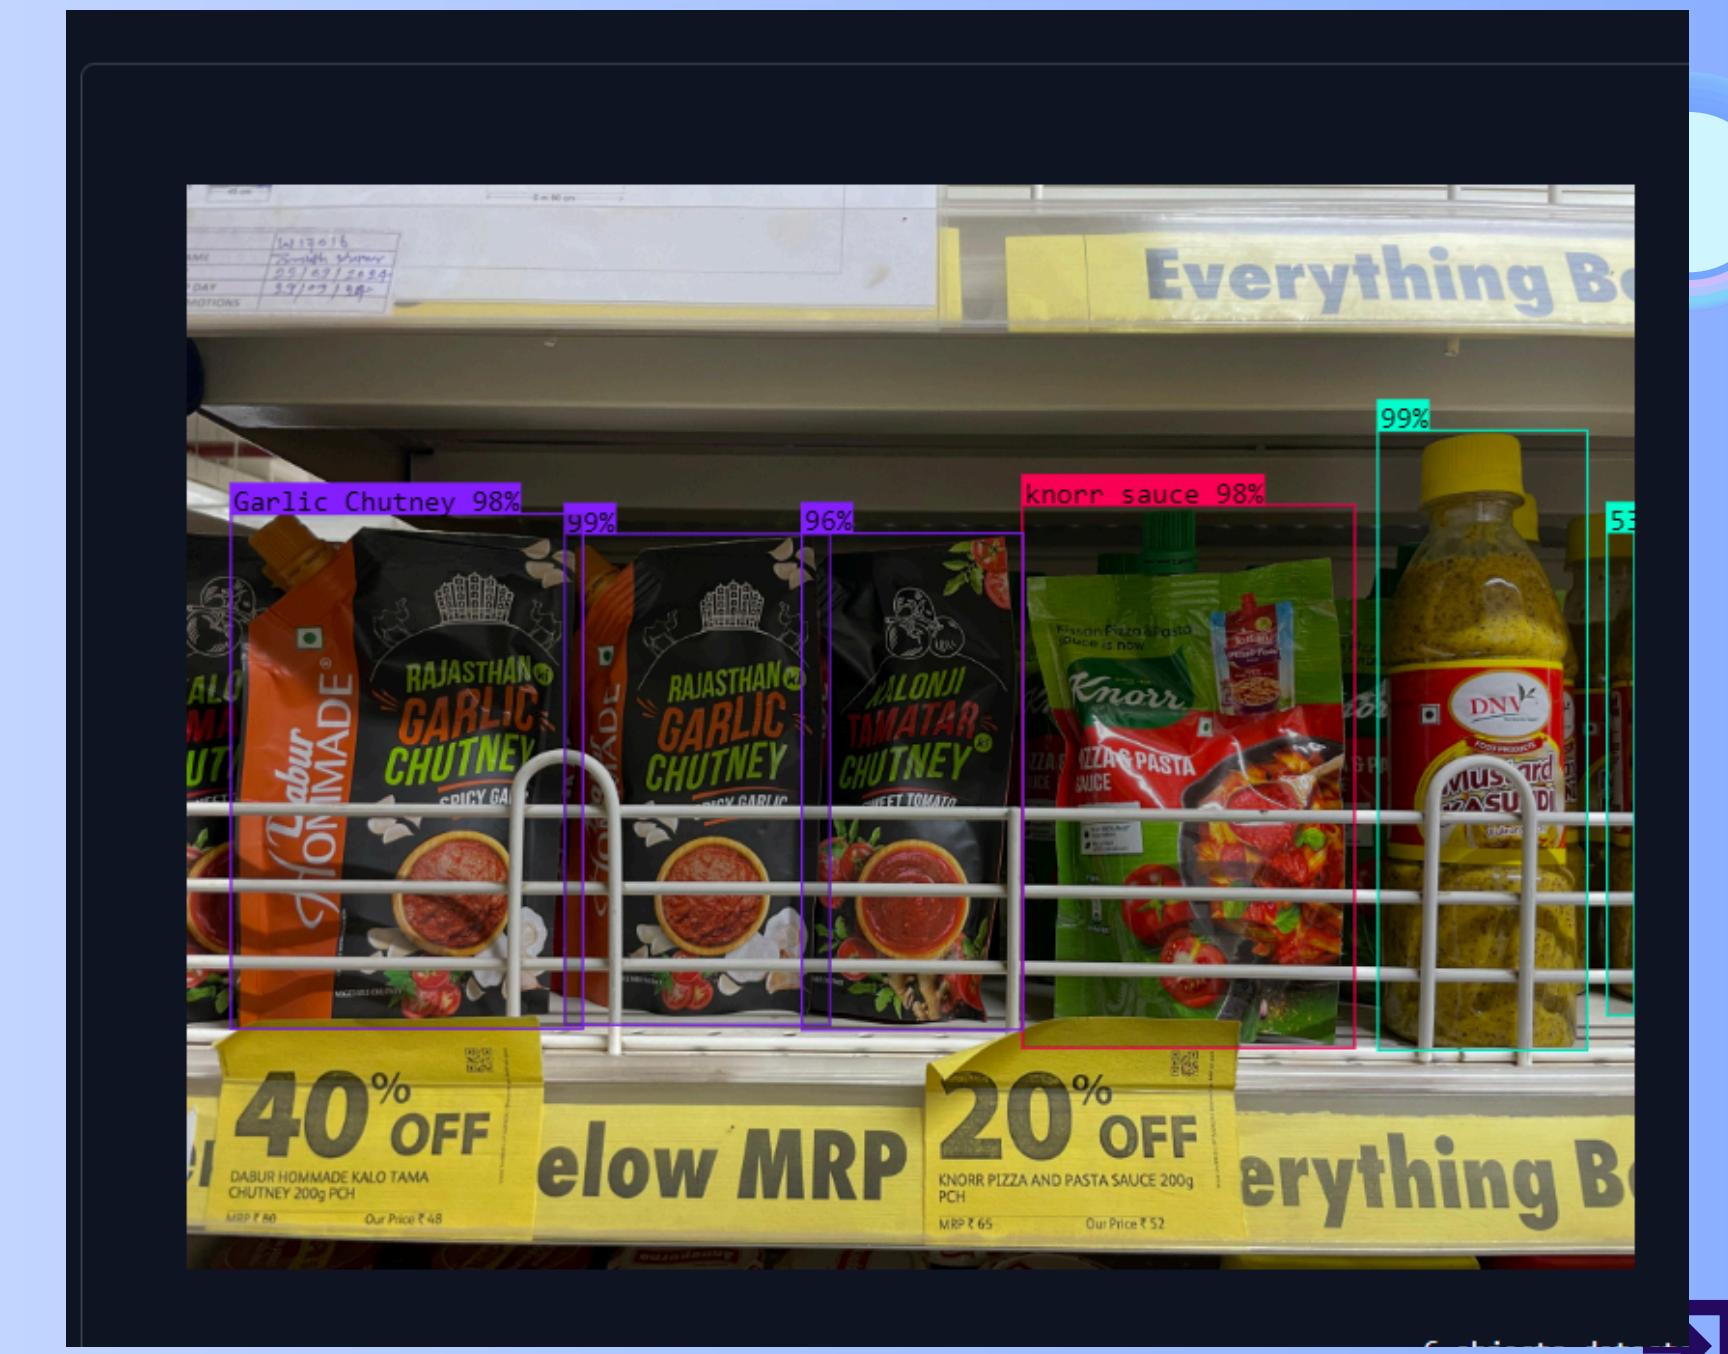
- RELEASED IN 2022, YOLOV7 IS OPTIMIZED FOR BOTH SPEED AND ACCURACY, MAKING IT ONE OF THE MOST EFFICIENT REAL-TIME OBJECT DETECTION MODELS ACHIEVES BETTER PERFORMANCE THAN PREVIOUS YOLO VERSIONS WITH FEWER PARAMETERS WHILE MAINTAINING HIGH ACCURACY.
- YOLOV7 IS DESIGNED WITH AN IMPROVED BACKBONE AND HEAD ARCHITECTURE, FOCUSING ON DETECTING SMALLER OBJECTS MORE EFFECTIVELY AND INCREASING DETECTION PRECISION.



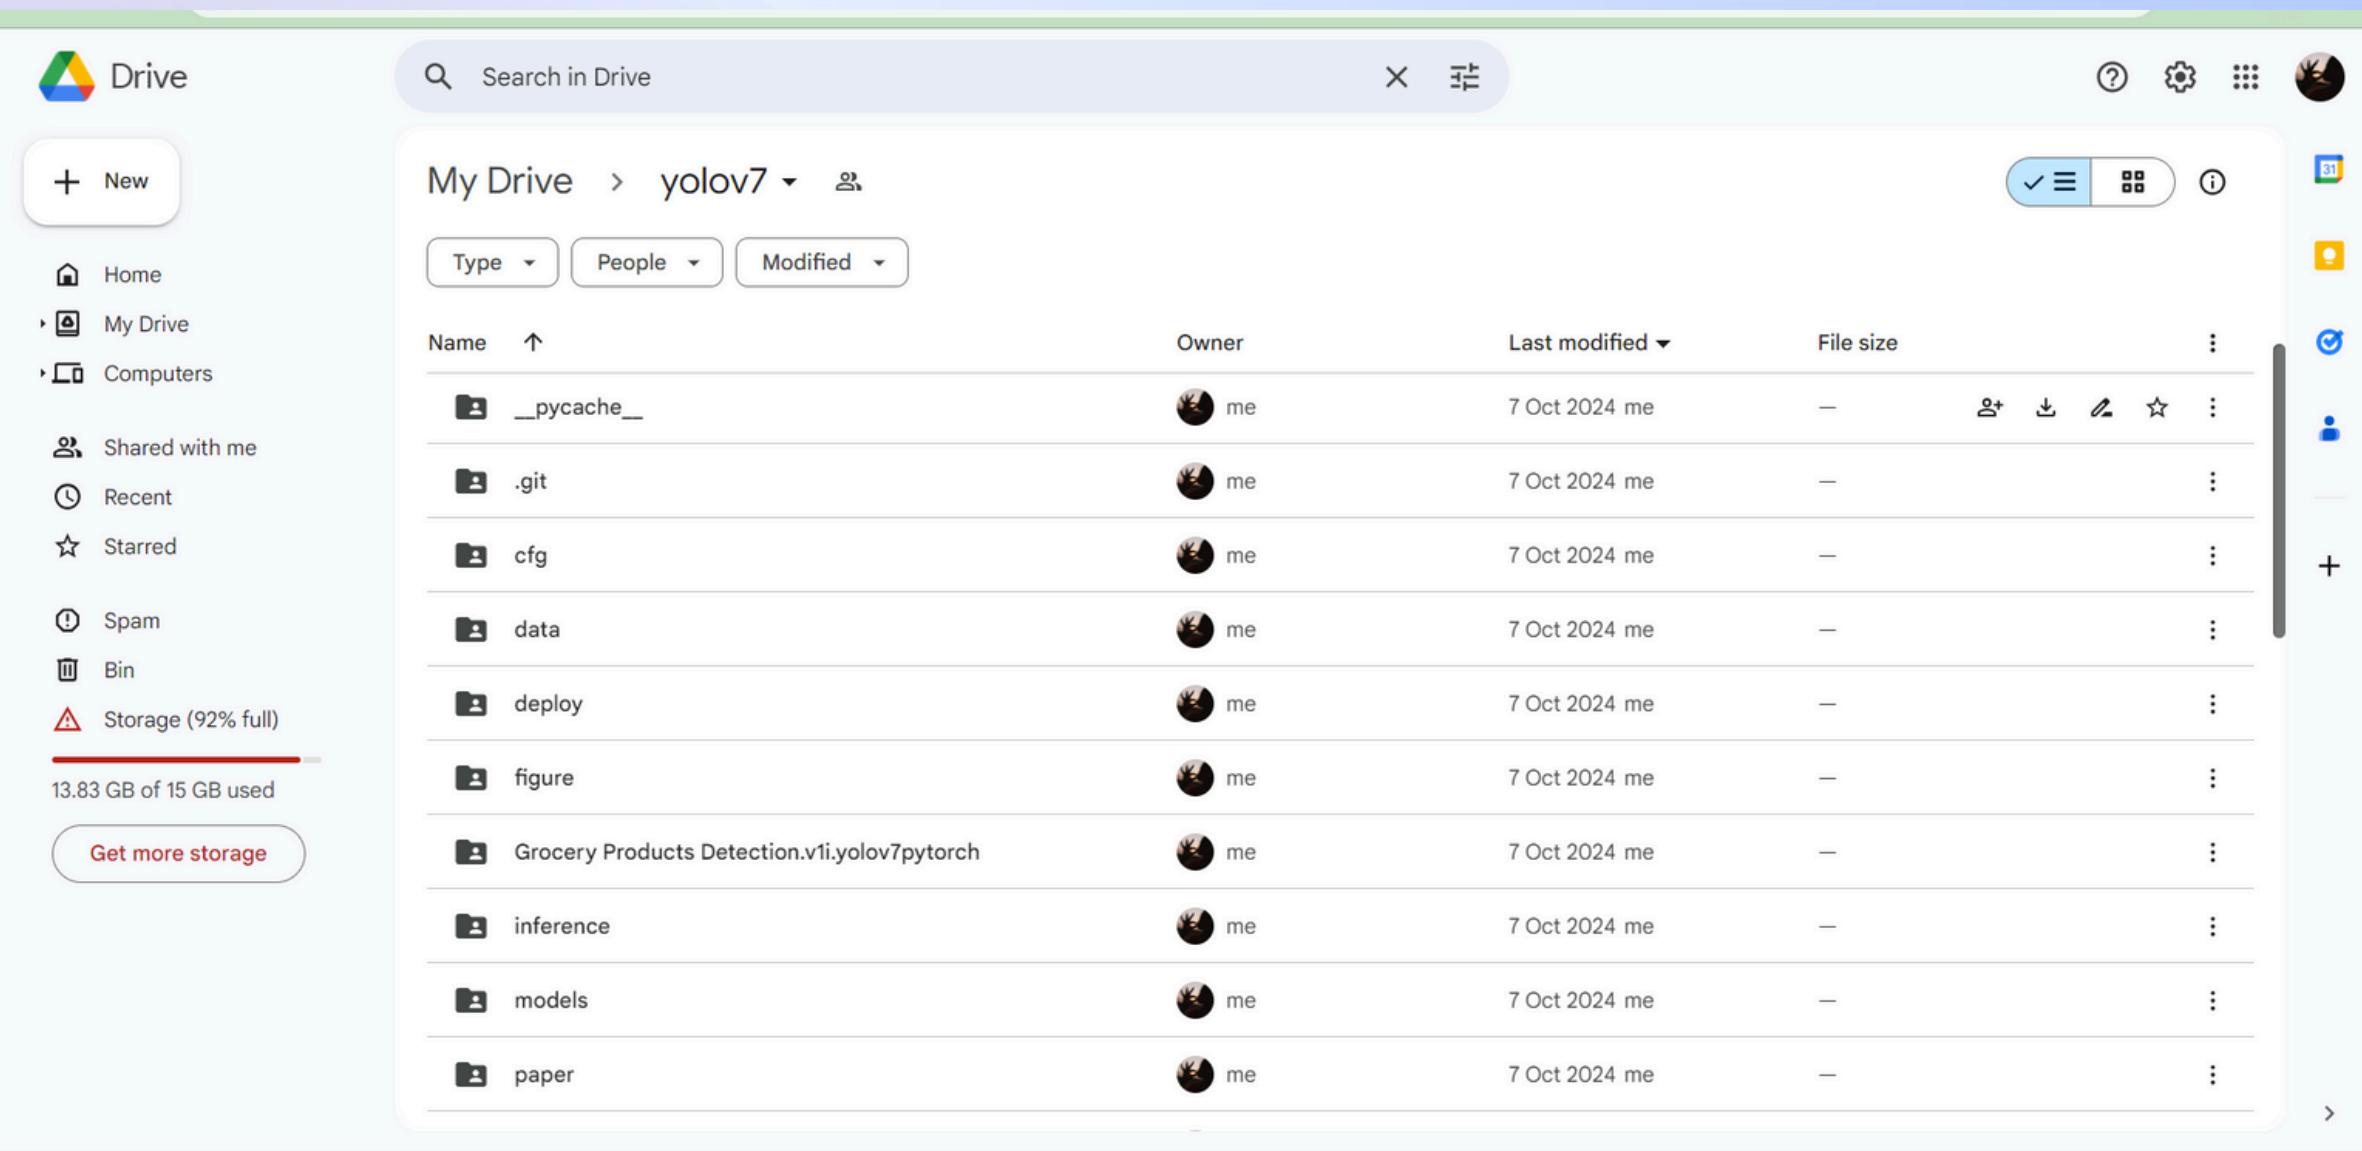
- LOSS FUNCTION USED IN YOLOV7:
  - BOUNDING BOX REGRESSION LOSS (CIOU): CALCULATES HOW WELL THE PREDICTED BOUNDING BOXES MATCH THE GROUND TRUTH.
  - OBJECTNESS LOSS: DETERMINES IF AN OBJECT IS PRESENT IN A GRID CELL.
  - CLASSIFICATION LOSS: HANDLES THE ACCURACY OF PREDICTING THE CORRECT CLASS LABEL FOR THE DETECTED OBJECTS.

# DATASET CREATION & ANNOTATION

- DATASET CREATED USING IMAGES OF SPECIFIC OBJECTS.
- ANNOTATION PLATFORM: ROBOFLOW USED FOR ANNOTATION.
- ANNOTATION INCLUDED DRAWING BOUNDING BOXES AROUND OBJECTS AND LABELING THEM



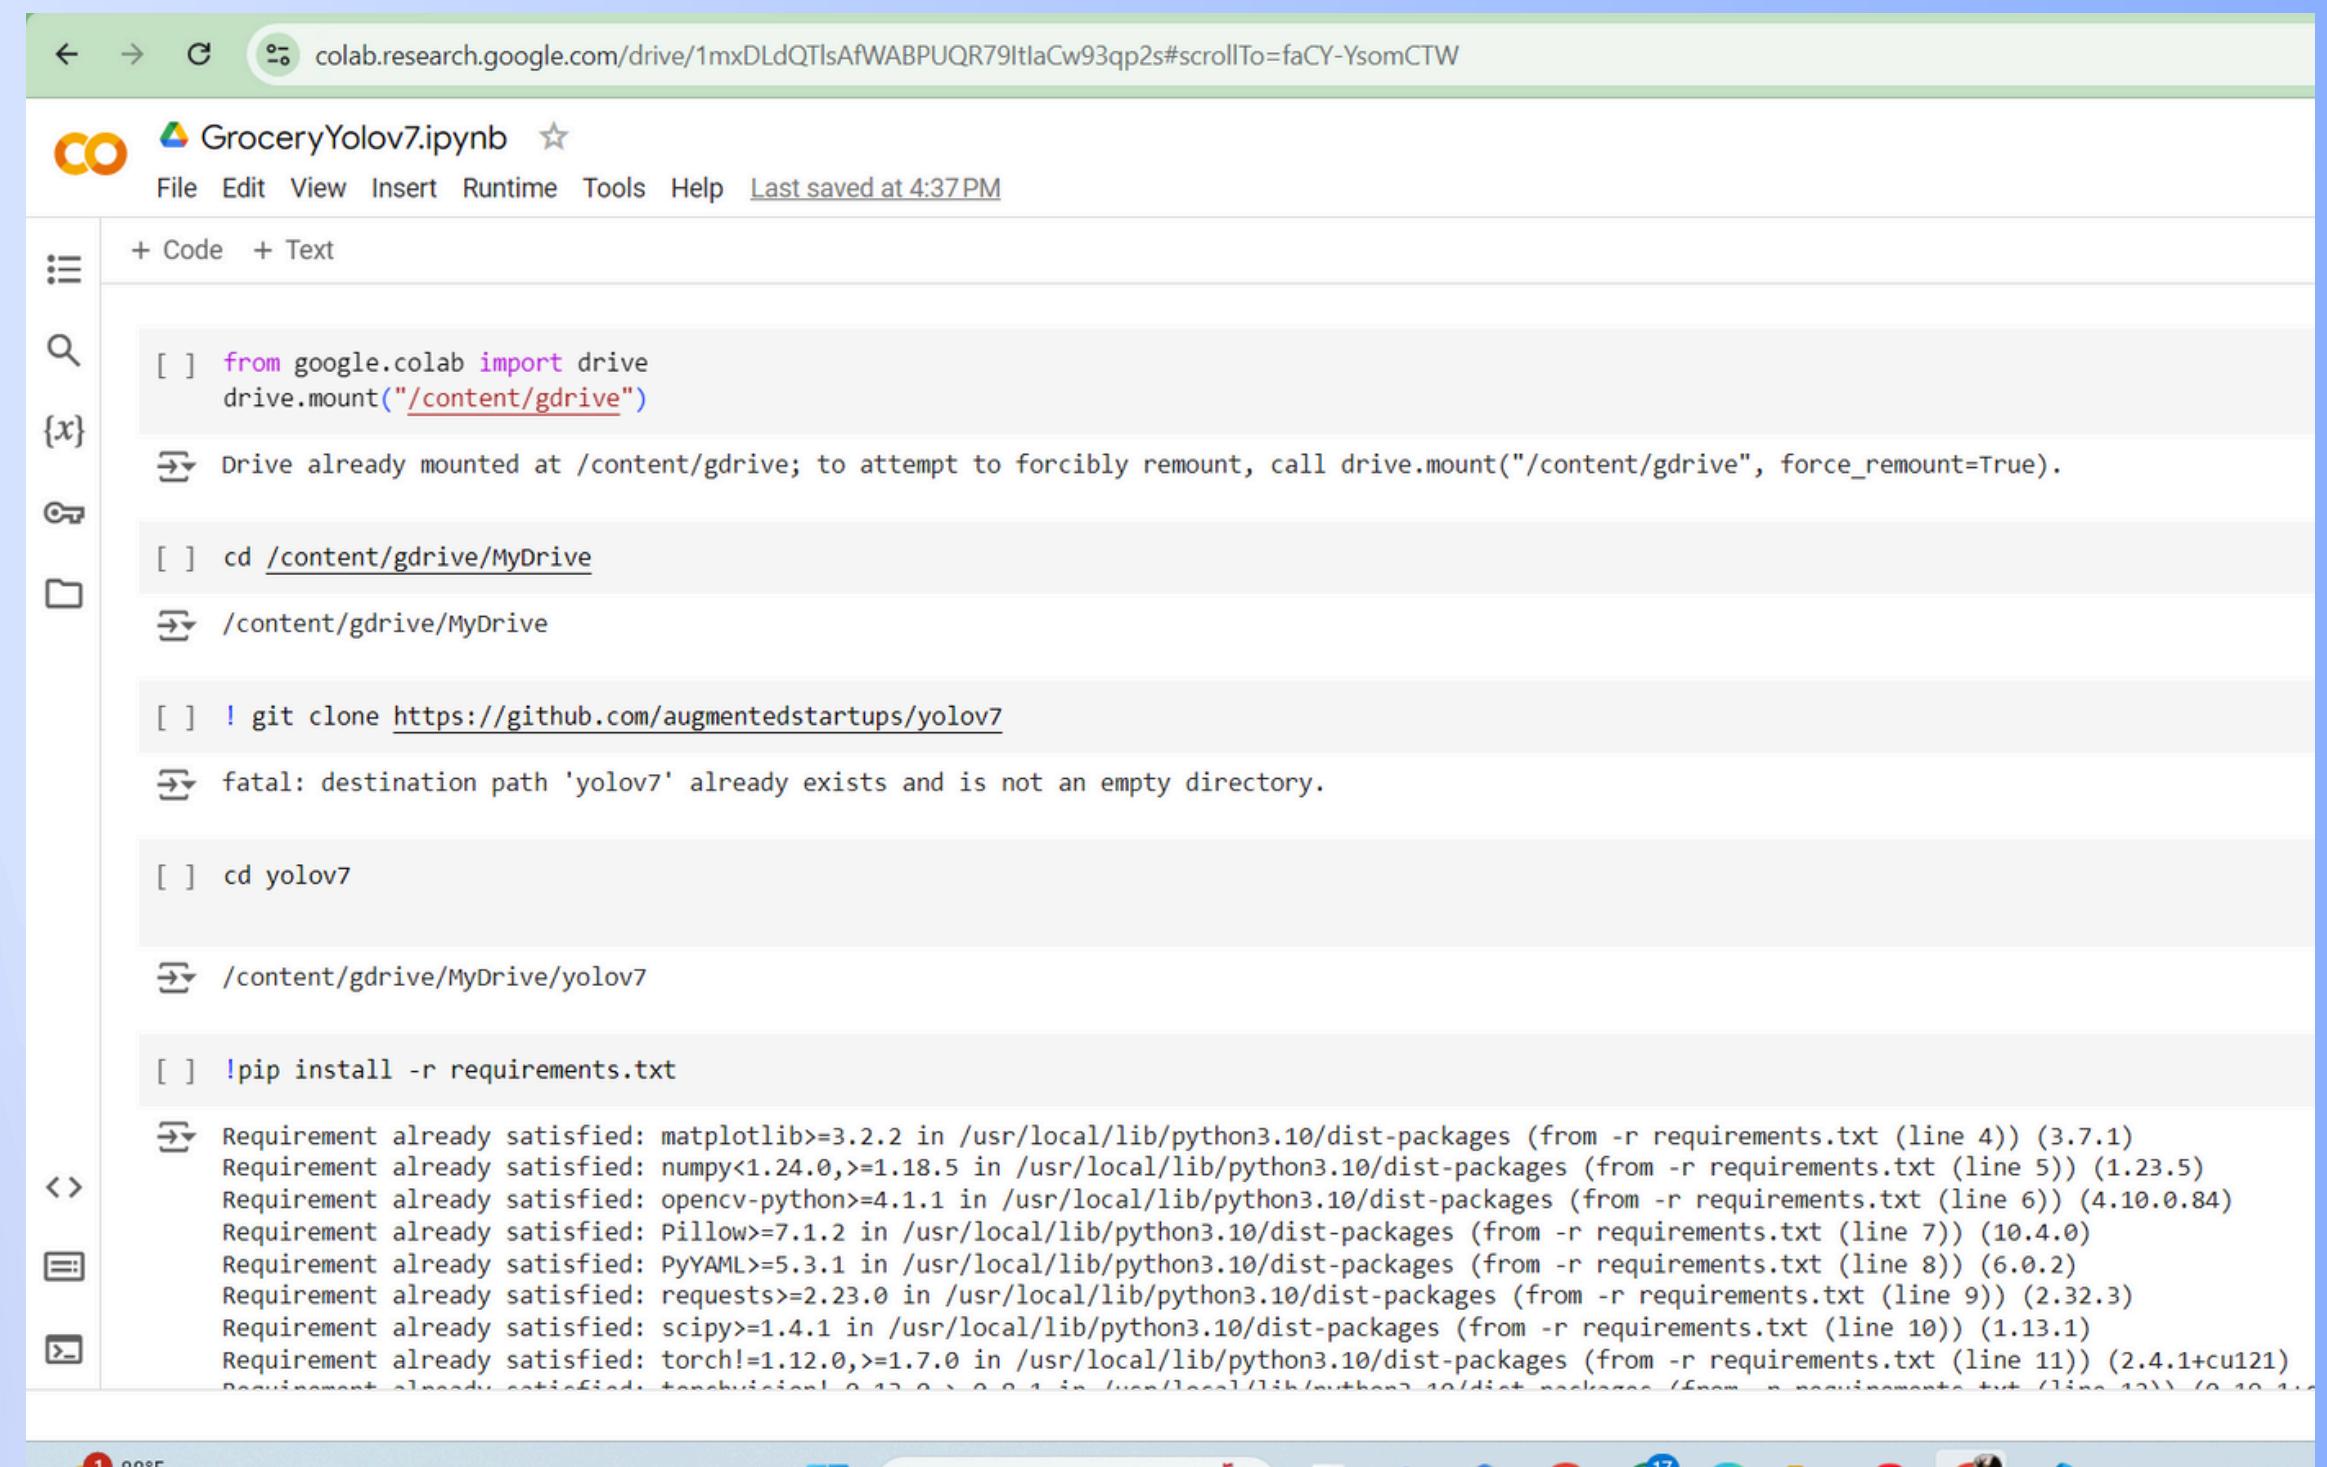
# TRAINING YOLOV7 MODEL



- **DATASET LOADED FROM ROBOFLOW, TO GOOGLE DRIVE WITH NECESSARY PRE-PROCESSING.**



- **PLATFORM: GOOGLE COLAB  
(LEVERAGE GPUS FOR FASTER TRAINING).**
- **FRAMEWORK: PYTORCH  
IMPLEMENTATION OF YOLOV7.**



The screenshot shows a Google Colab notebook titled "GroceryYolov7.ipynb". The code cell contains the following commands:

```
[ ] from google.colab import drive
drive.mount("/content/gdrive")
[ ] cd /content/gdrive/MyDrive
[ ] ! git clone https://github.com/augmentedstartups/yolov7
[ ] cd yolov7
[ ] ! pip install -r requirements.txt
```

Output from the first command:

```
→ Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).
```

Output from the second command:

```
→ /content/gdrive/MyDrive
```

Output from the third command:

```
→ fatal: destination path 'yolov7' already exists and is not an empty directory.
```

Output from the fourth command:

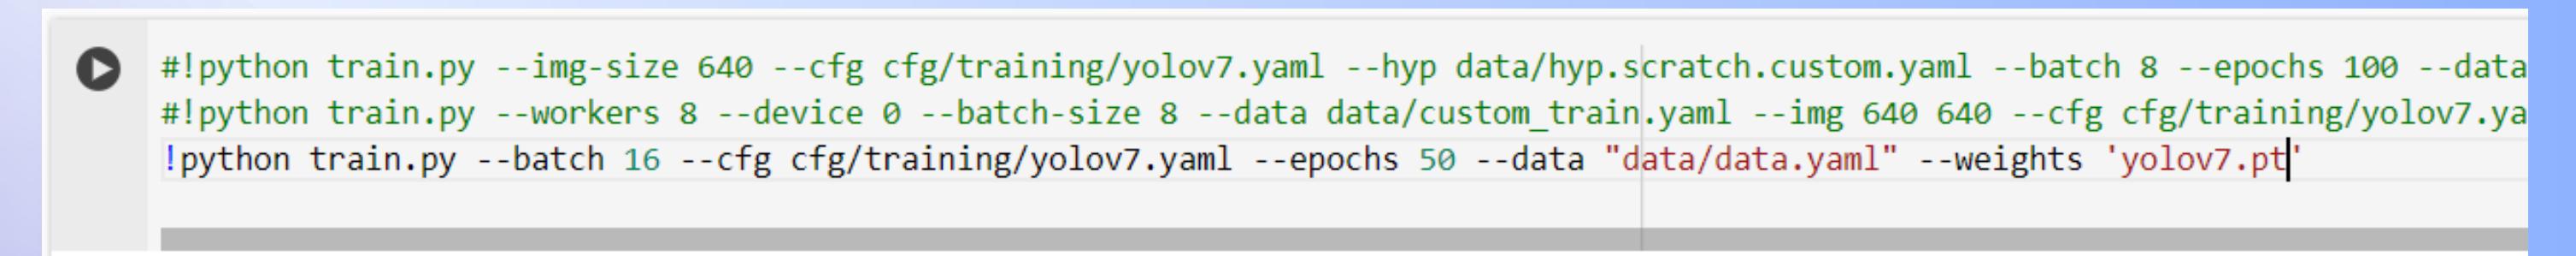
```
→ /content/gdrive/MyDrive/yolov7
```

Output from the fifth command:

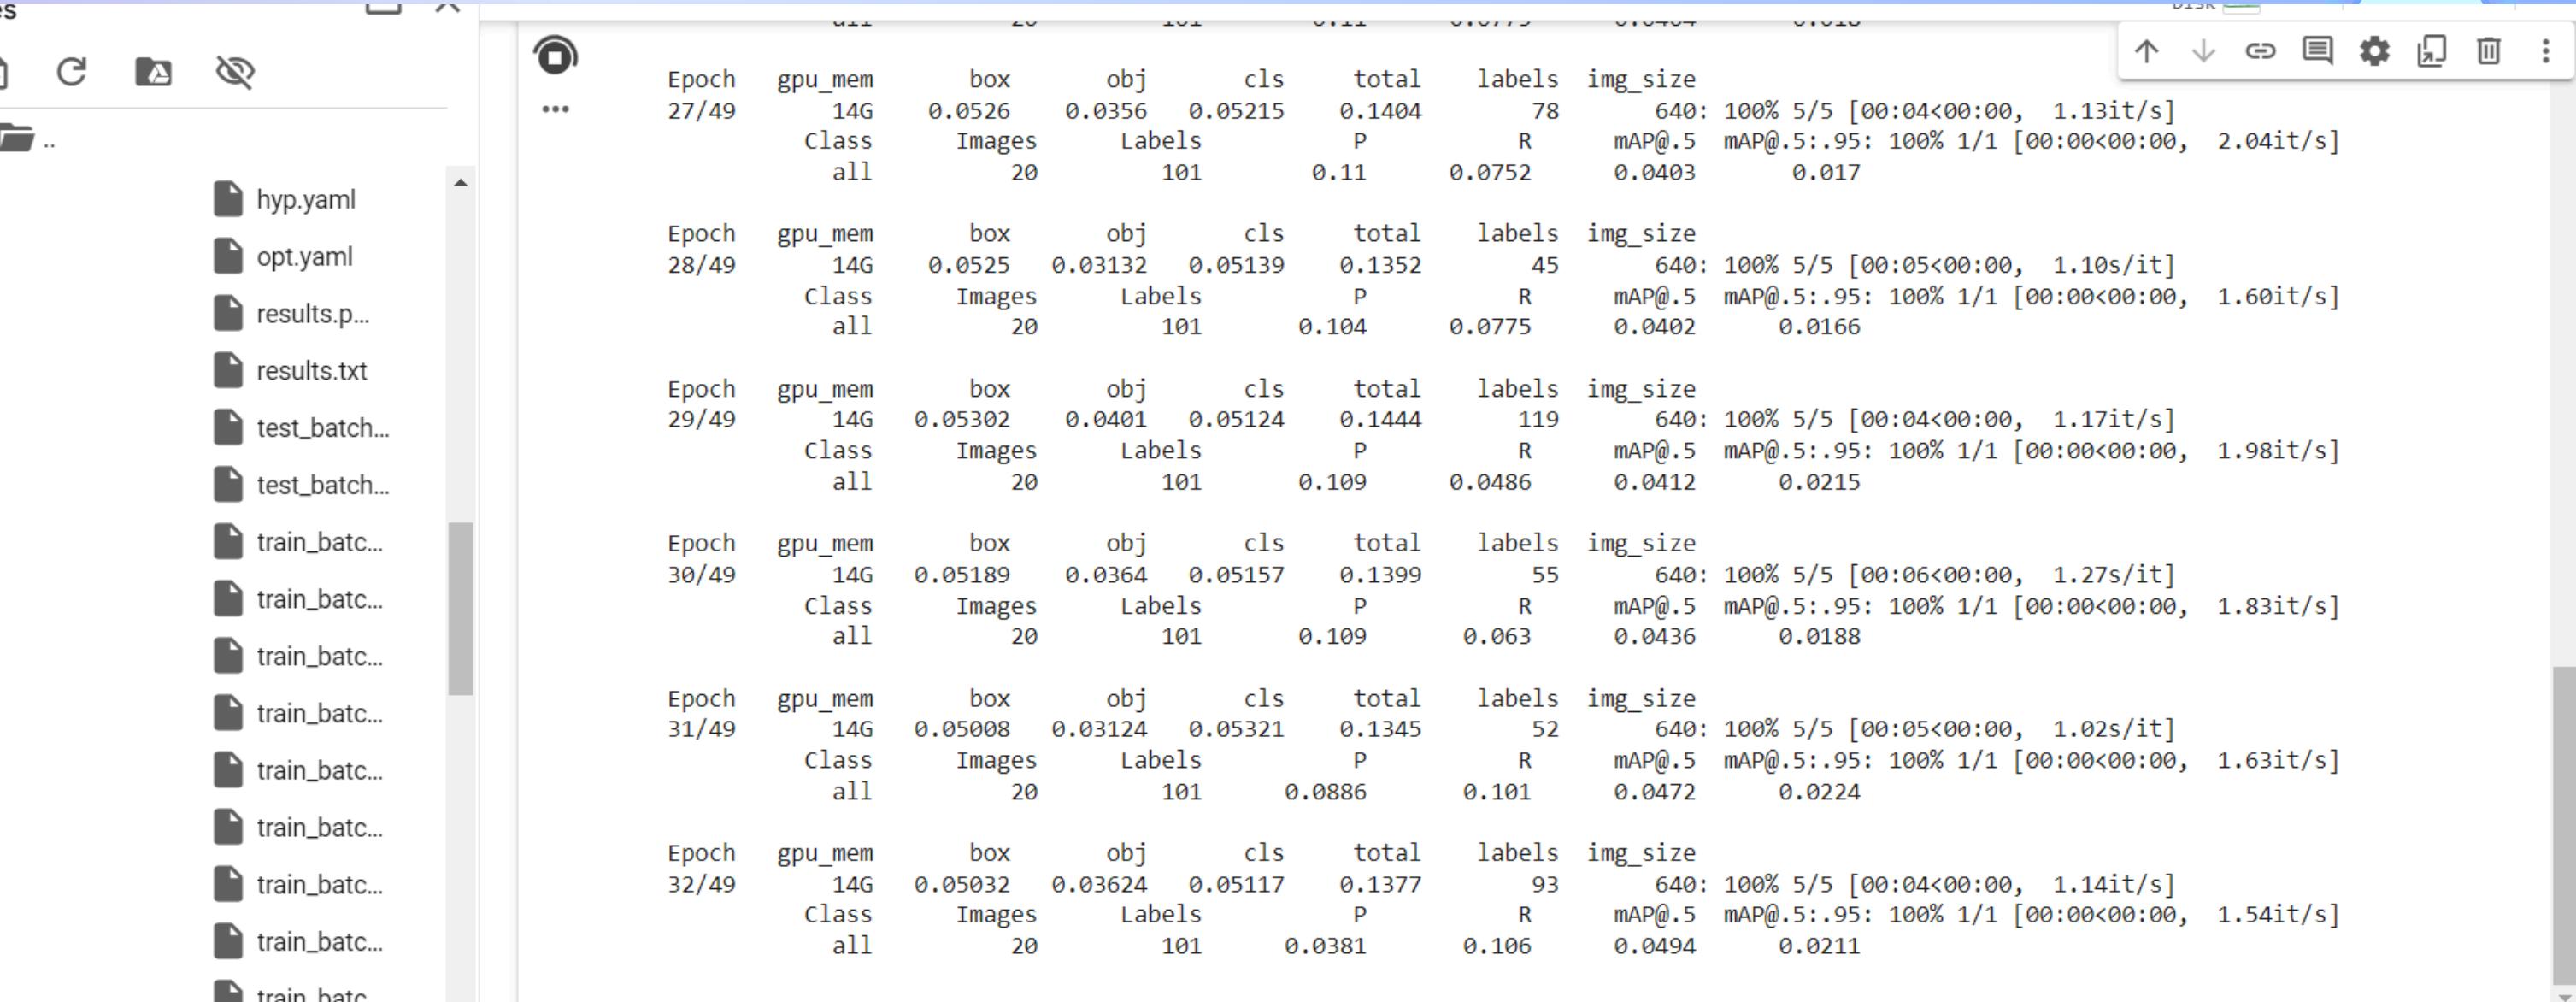
```
→ Requirement already satisfied: matplotlib>=3.2.2 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 4)) (3.7.1)
Requirement already satisfied: numpy<1.24.0,>=1.18.5 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 5)) (1.23.5)
Requirement already satisfied: opencv-python>=4.1.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 6)) (4.10.0.84)
Requirement already satisfied: Pillow>=7.1.2 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 7)) (10.4.0)
Requirement already satisfied: PyYAML>=5.3.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 8)) (6.0.2)
Requirement already satisfied: requests>=2.23.0 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 9)) (2.32.3)
Requirement already satisfied: scipy>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 10)) (1.13.1)
Requirement already satisfied: torch!=1.12.0,>=1.7.0 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 11)) (2.4.1+cu121)
Requirement already satisfied: torchvision<0.12.0,>=0.9.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 12)) (0.10.0)
```

## • STARTED THE TRAINING PROCESS WITH PRE-TRAINED WEIGHTS

[TTPS://GITHUB.COM/WONGKINYIU/YOLOV7/RELEASES/DOWNLOAD/V0.1/YOLOV7.PT](https://github.com/wongkinyiu/yolov7/releases/download/v0.1/yolov7.pt)



```
#!python train.py --img-size 640 --cfg cfg/training/yolov7.yaml --hyp data/hyp.scratch.custom.yaml --batch 8 --epochs 100 --data data/custom_train.yaml  
#!python train.py --workers 8 --device 0 --batch-size 8 --data data/custom_train.yaml --img 640 640 --cfg cfg/training/yolov7.yaml  
!python train.py --batch 16 --cfg cfg/training/yolov7.yaml --epochs 50 --data "data/data.yaml" --weights 'yolov7.pt'
```



Epoch	gpu_mem	box	obj	cls	total	labels	img_size
27/49	14G	0.0526	0.0356	0.05215	0.1404	78	640: 100% 5/5 [00:04<00:00, 1.13it/s]
	Class Images		Labels		P	R	MAP@.5 MAP@.5:.95: 100% 1/1 [00:00<00:00, 2.04it/s]
	all	20		101	0.11	0.0752	0.0403 0.017
28/49	14G	0.0525	0.03132	0.05139	0.1352	45	640: 100% 5/5 [00:05<00:00, 1.10s/it]
	Class Images		Labels		P	R	MAP@.5 MAP@.5:.95: 100% 1/1 [00:00<00:00, 1.60it/s]
	all	20		101	0.104	0.0775	0.0402 0.0166
29/49	14G	0.05302	0.0401	0.05124	0.1444	119	640: 100% 5/5 [00:04<00:00, 1.17it/s]
	Class Images		Labels		P	R	MAP@.5 MAP@.5:.95: 100% 1/1 [00:00<00:00, 1.98it/s]
	all	20		101	0.109	0.0486	0.0412 0.0215
30/49	14G	0.05189	0.0364	0.05157	0.1399	55	640: 100% 5/5 [00:06<00:00, 1.27s/it]
	Class Images		Labels		P	R	MAP@.5 MAP@.5:.95: 100% 1/1 [00:00<00:00, 1.83it/s]
	all	20		101	0.109	0.063	0.0436 0.0188
31/49	14G	0.05008	0.03124	0.05321	0.1345	52	640: 100% 5/5 [00:05<00:00, 1.02s/it]
	Class Images		Labels		P	R	MAP@.5 MAP@.5:.95: 100% 1/1 [00:00<00:00, 1.63it/s]
	all	20		101	0.0886	0.101	0.0472 0.0224
32/49	14G	0.05032	0.03624	0.05117	0.1377	93	640: 100% 5/5 [00:04<00:00, 1.14it/s]
	Class Images		Labels		P	R	MAP@.5 MAP@.5:.95: 100% 1/1 [00:00<00:00, 1.54it/s]
	all	20		101	0.0381	0.106	0.0494 0.0211

# • AFTER THE FIRST TRAINING PROCESS WE GOT THE UPDATED WEIGHTS

slov7.ipynb ★

Insert Runtime Tools Help All changes saved

+ Code + Text

9m

Class all Images 20 Labels 101 P 0.0253 R 0.315 mAP@.5 mAP@.5:.95: 100% 1/1 [00:00<0:00:00] 0.0631 0.0325

Epoch 45/49 gpu\_mem box obj cls total labels img\_size 14G 0.04759 0.04035 0.05139 0.1393 115 640: 100% 5/5 [00:04<00:00, 1.03it/s]

Class all Images 20 Labels 101 P 0.0265 R 0.201 mAP@.5 mAP@.5:.95: 100% 1/1 [00:00<00:00, 1.58it/s] 0.0595 0.0325

Epoch 46/49 gpu\_mem box obj cls total labels img\_size 14G 0.04699 0.03576 0.05037 0.1331 145 640: 100% 5/5 [00:04<00:00, 1.15it/s]

Class all Images 20 Labels 101 P 0.0255 R 0.451 mAP@.5 mAP@.5:.95: 100% 1/1 [00:00<00:00, 2.19it/s] 0.0634 0.0314

Epoch 47/49 gpu\_mem box obj cls total labels img\_size 14G 0.04669 0.03969 0.05107 0.1375 114 640: 100% 5/5 [00:05<00:00, 1.15s/it]

Class all Images 20 Labels 101 P 0.0257 R 0.451 mAP@.5 mAP@.5:.95: 100% 1/1 [00:01<00:00, 1.06s/it] 0.0631 0.0336

Epoch 48/49 gpu\_mem box obj cls total labels img\_size 14G 0.04691 0.03625 0.05127 0.1344 61 640: 100% 5/5 [00:04<00:00, 1.12it/s]

Class all Images 20 Labels 101 P 0.029 R 0.23 mAP@.5 mAP@.5:.95: 100% 1/1 [00:00<00:00, 1.99it/s] 0.0643 0.0342

Epoch 49/49 gpu\_mem box obj cls total labels img\_size 14G 0.04651 0.03642 0.05049 0.1334 81 640: 100% 5/5 [00:05<00:00, 1.16s/it]

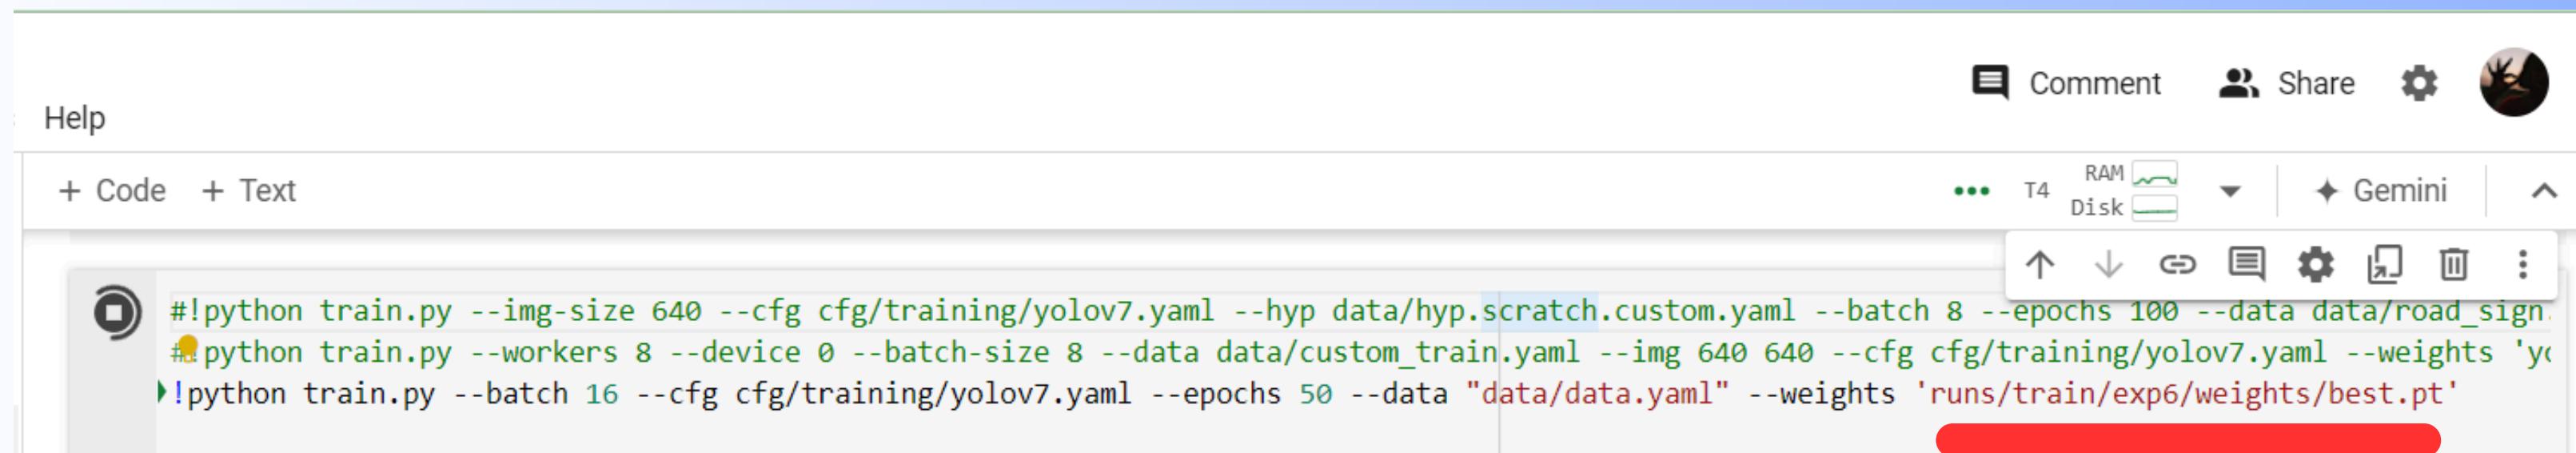
Class all Images 20 Labels 101 P 0.0309 R 0.234 mAP@.5 mAP@.5:.95: 100% 1/1 [00:01<00:00, 1.23s/it] 0.0653 0.0354

50 epochs completed in 0.147 hours.

/content/gdrive/MyDrive/yolov7/utils/general.py:802: FutureWarning: You are using `torch.load` with `weights\_only=False` (the current default) x = torch.load(f, map\_location=torch.device('cpu')) Optimizer stripped from runs/train/exp6/weights/last.pt, 75.3MB Optimizer stripped from runs/train/exp6/weights/best.pt, 75.3MB

71.69 GB available

- BY REPLACING THE OLD WEIGHTS WITH UPDATED WEIGHTS WE HAVE TRAINED THE MODEL MULTIPLE TIMES TO INCREASE THE MEAN AVERAGE PRECISION



A screenshot of a Jupyter Notebook cell. The cell contains three lines of code:

```
#!python train.py --img-size 640 --cfg cfg/training/yolov7.yaml --hyp data/hyp.scratch.custom.yaml --batch 8 --epochs 100 --data data/road_signs
#!python train.py --workers 8 --device 0 --batch-size 8 --data data/custom_train.yaml --img 640 640 --cfg cfg/training/yolov7.yaml --weights 'yolov7.pt'
!python train.py --batch 16 --cfg cfg/training/yolov7.yaml --epochs 50 --data "data/data.yaml" --weights 'runs/train/exp6/weights/best.pt'
```

- **UPDATED WEIGHTS**

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
48/49	14.6G	0.03903	0.03722	0.04958	0.1258	61	640: 100% 5/5 [00:05<00:00, 1.04s/it]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 1/1 [00:01<00:00, 1.11s/it]
	all	20	101	0.0456	0.341	0.0855	0.0576

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
49/49	14.6G	0.03828	0.03702	0.04875	0.1241	81	640: 100% 5/5 [00:04<00:00, 1.14it/s]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 1/1 [00:01<00:00, 1.43s/it]
	all	20	101	0.239	0.184	0.0881	0.0613

50 epochs completed in 0.150 hours.

```
/content/gdrive/MyDrive/yolov7/utils/general.py:802: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default)
  x = torch.load(f, map_location=torch.device('cpu'))
Optimizer stripped from runs/train/exp7/weights/last.pt, 75.3MB
Optimizer stripped from runs/train/exp7/weights/best.pt, 75.3MB
```

## • TRAINING WITH UPDATED WEIGHTS

```
#!python train.py --img-size 640 --cfg cfg/training/yolov7.yaml --hyp data/hyp.scratch.custom.yaml --batch 8 --epochs 100 --data data/road_signs
#!python train.py --workers 8 --device 0 --batch-size 8 --data data/custom_train.yaml --img 640 640 --cfg cfg/training/yolov7.yaml --weights 'yolov7.pt'
!python train.py --batch 16 --cfg cfg/training/yolov7.yaml --epochs 50 --data "data/data.yaml" --weights 'runs/train/exp11/weights/best.pt'

...
all      20      101     0.0598      0.426      0.0995      0.0724

Epoch  gpu_mem      box      obj      cls      total      labels      img_size
34/49    15.1G    0.0362    0.03474   0.04856   0.1195      85        640: 100% 5/5 [00:06<00:00, 1.20s/it]
          Class      Images     Labels           P           R      mAP@.5  mAP@.5:.95: 100% 1/1 [00:00<00:00, 1.06it/s]
          all       20       101     0.0586      0.44      0.0957      0.0676

Epoch  gpu_mem      box      obj      cls      total      labels      img_size
35/49    15.1G    0.03682   0.03022   0.04729   0.1143      35        640: 100% 5/5 [00:04<00:00, 1.00it/s]
          Class      Images     Labels           P           R      mAP@.5  mAP@.5:.95: 100% 1/1 [00:00<00:00, 1.09it/s]
```

## • NEW UPDATED WEIGHTS

```
Epoch  gpu_mem      box      obj      cls      total      labels      img_size
48/49    15.1G    0.03422   0.03678   0.04826   0.1193      61        640: 100% 5/5 [00:05<00:00, 1.02s/it]
          Class      Images     Labels           P           R      mAP@.5  mAP@.5:.95: 100% 1/1 [00:00<00:00, 1.86it/s]
          all       20       101     0.0582      0.329      0.105      0.0819

Epoch  gpu_mem      box      obj      cls      total      labels      img_size
49/49    15.1G    0.03361   0.03636   0.0472    0.1172      81        640: 100% 5/5 [00:06<00:00, 1.36s/it]
          Class      Images     Labels           P           R      mAP@.5  mAP@.5:.95: 100% 1/1 [00:01<00:00, 1.25s/it]
          all       20       101     0.06      0.325      0.105      0.077

50 epochs completed in 0.149 hours.
```

```
/content/gdrive/MyDrive/yolov7/utils/general.py:802: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default)
  x = torch.load(f, map_location=torch.device('cpu'))
Optimizer stripped from runs/train/exp12/weights/last.pt, 75.3MB
Optimizer stripped from runs/train/exp12/weights/best.pt, 75.3MB
```

## • USING THE FINAL WEIGHTS WE HAVE PREDICTED OUTPUTS FOR OUR IMAGES

```
[x] !python detect.py --weights '/content/gdrive/MyDrive/yolov7/runs/train/exp13/weights/last.pt' --conf 0.1 --source '/content/gdrive/MyDrive/yolov7/Grocery Products Detection.v1i.yolov7pytorch/test/images'  
Namespace(weights=['/content/gdrive/MyDrive/yolov7/runs/train/exp13/weights/last.pt'], source='/content/gdrive/MyDrive/yolov7/Grocery Products Detection.v1i.yolov7pytorch/test/images',  
YOLOR 🚀 v0.1-121-g2fdc7f1 torch 2.4.1+cu121 CUDA:0 (Tesla T4, 15102.0625MB)  
  
/content/gdrive/MyDrive/yolov7/models/experimental.py:252: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the default pickl  
ckpt = torch.load(w, map_location=map_location) # load  
Fusing layers...  
RepConv.fuse_repvgg_block  
RepConv.fuse_repvgg_block  
RepConv.fuse_repvgg_block  
IDetect.fuse  
/usr/local/lib/python3.10/dist-packages/tor  
✓ T4 RAM Disk Gemini ↑  
IMG20241003103119.jpg.rf.4cc03dc16f586a2606fad1 be required to pass the indexing argument. (Triggered inter
```

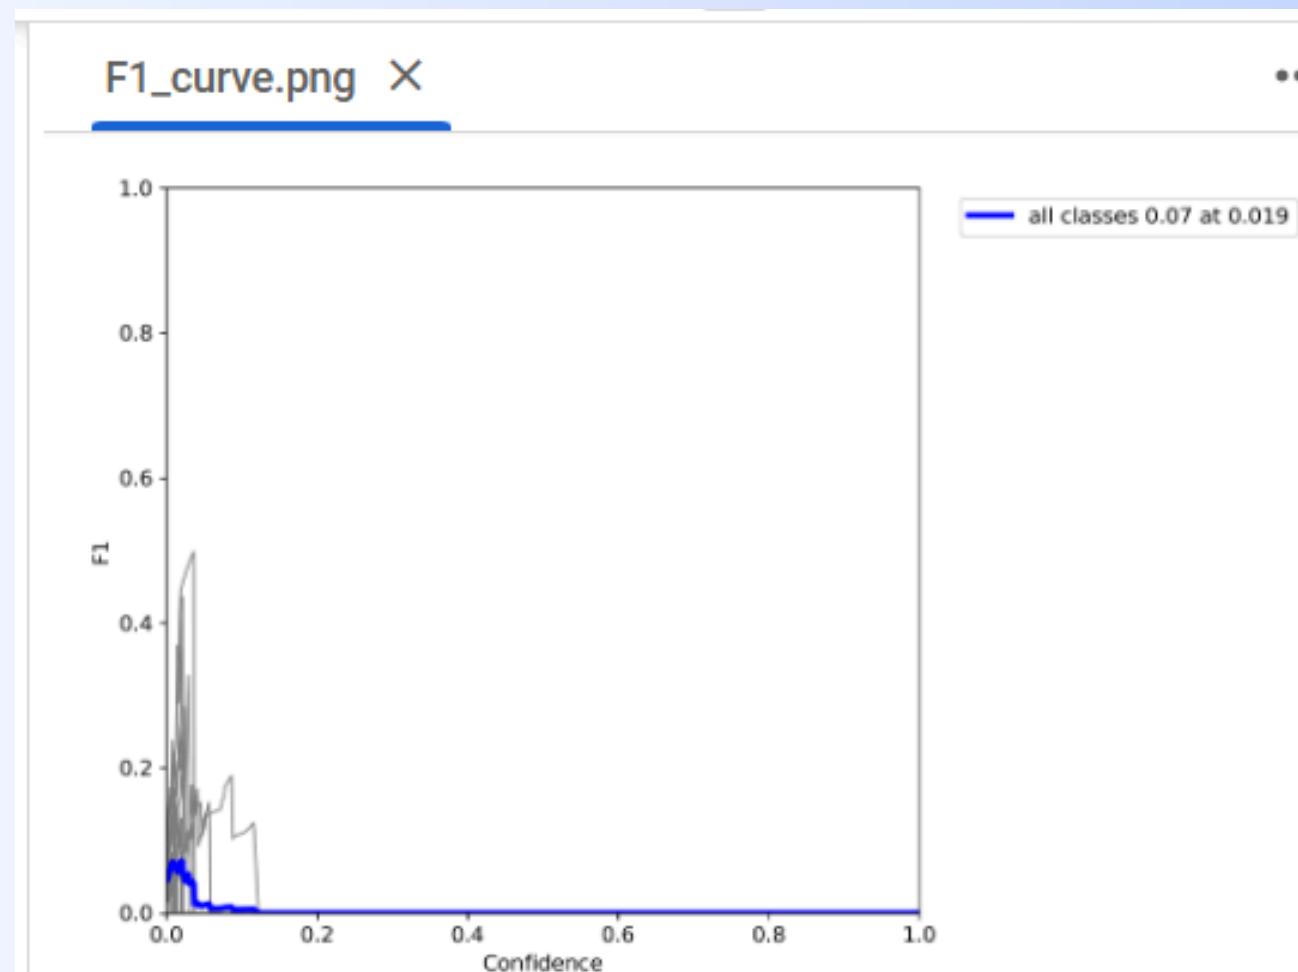


# MODEL PERFORMANCE AND EVALUATION

- **MAP (MEAN AVERAGE PRECISION): MEASURES THE ACCURACY OF THE MODEL IN DETECTING AND CLASSIFYING OBJECTS.**

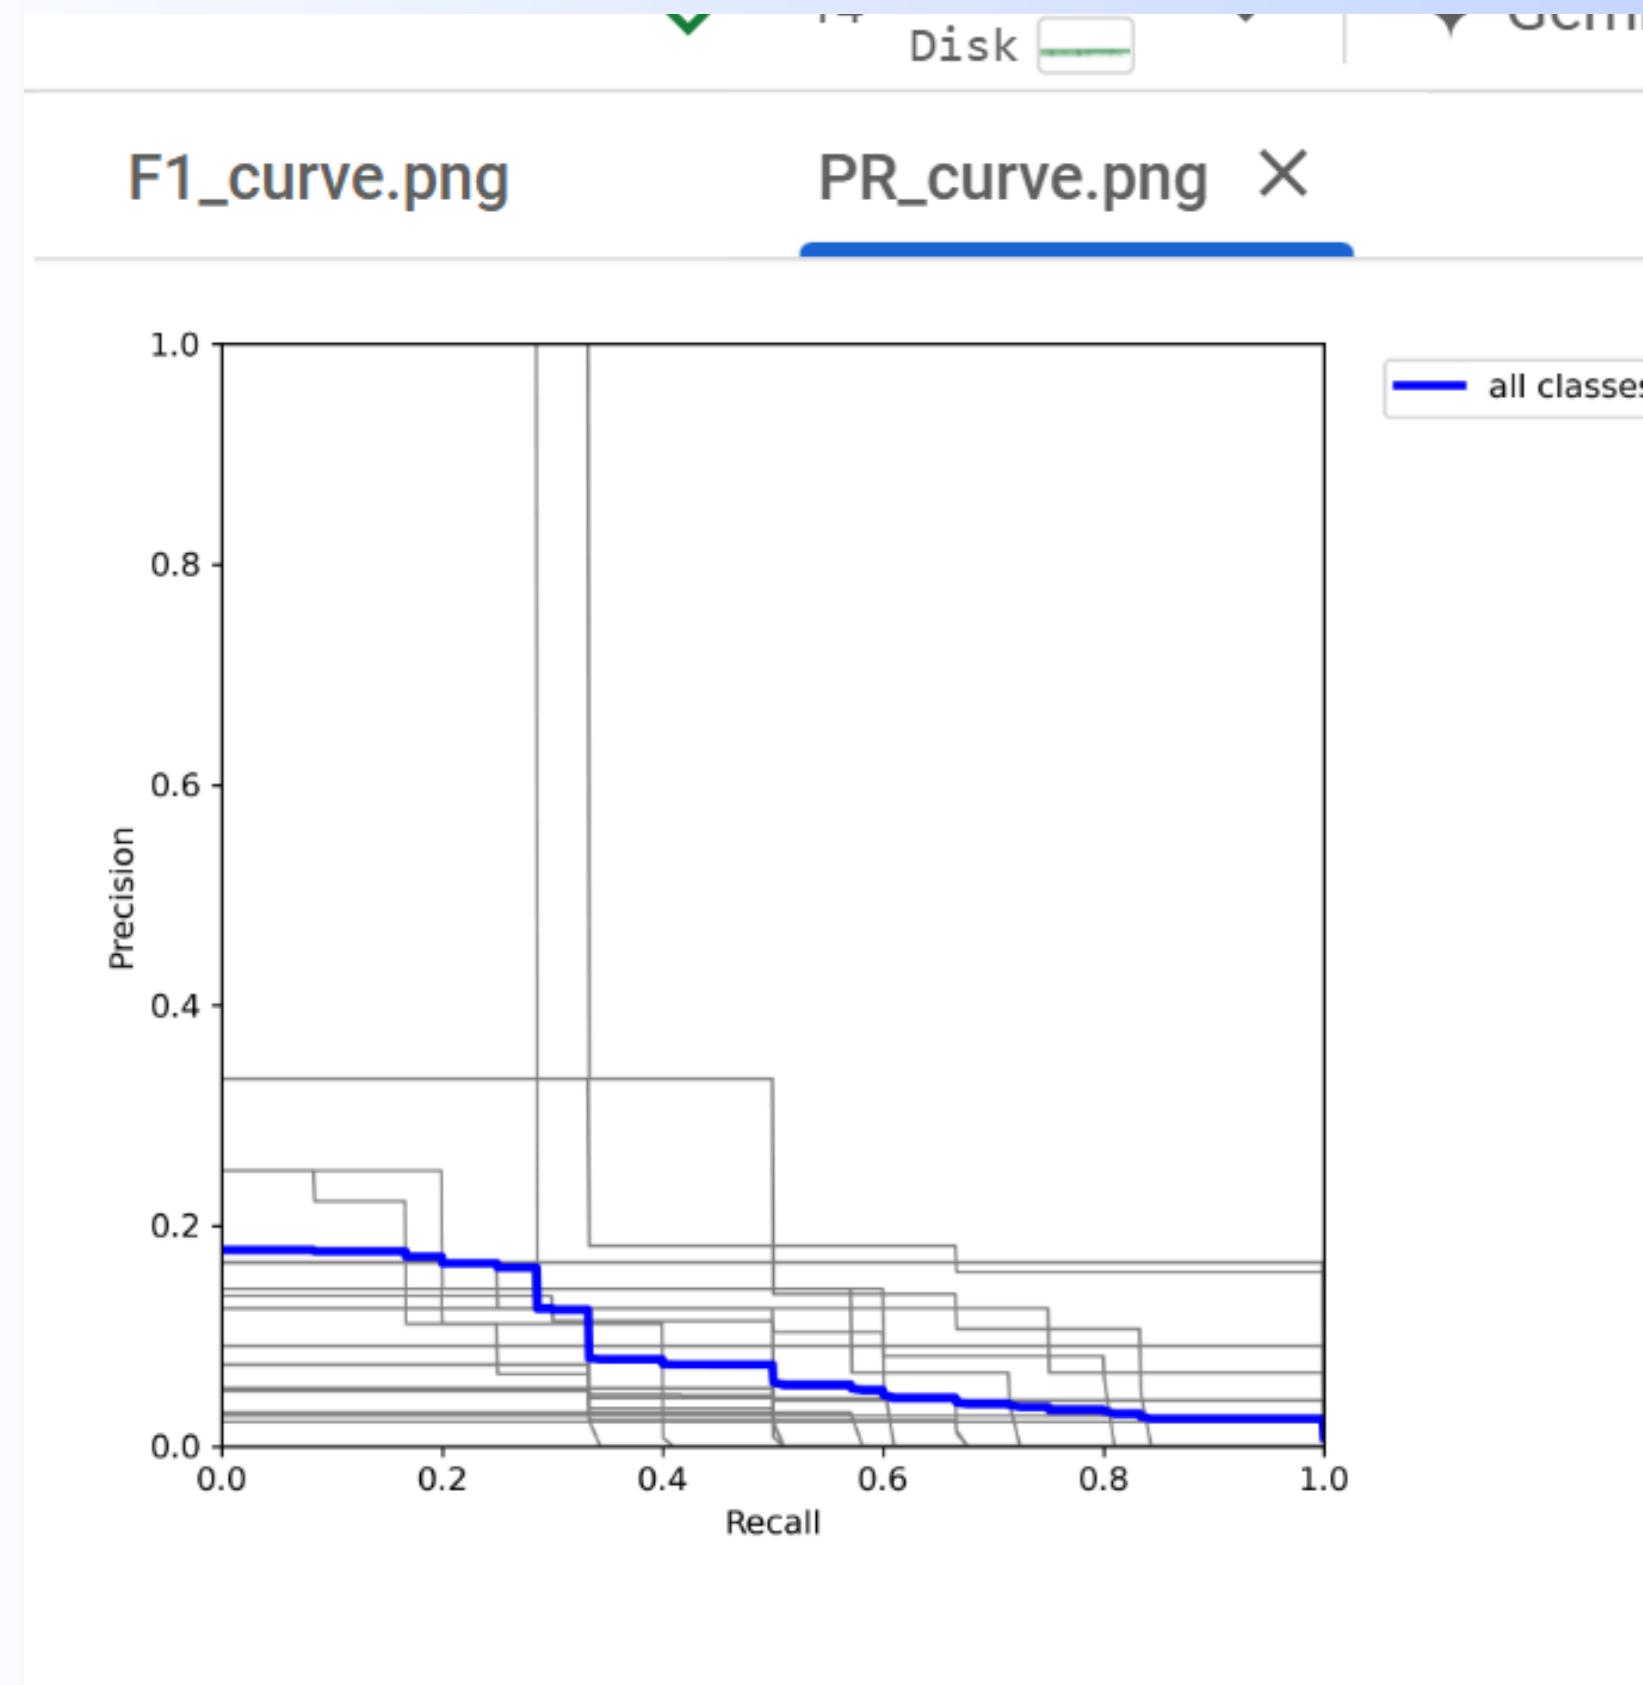
P	R	mAP@.5
0.222	0.327	0.115

- **F1 CURVE**



**OBSERVATION:**  
**F1-SCORE ACROSS ALL CLASSES IS 0.07**  
**WHEN THE CONFIDENCE THRESHOLD IS SET**  
**AT 0.019. THIS VALUE SEEMS LOW,**  
**WHICH MAY INDICATE THAT YOUR MODEL**  
**IS NOT PERFORMING WELL IN TERMS OF**  
**BALANCING PRECISION AND RECALL.**

## • PRECISION-RECALL CURVE



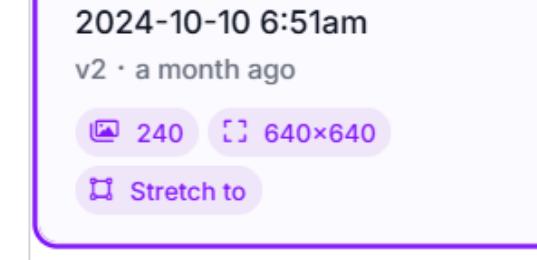
**OBSERVATION:**  
FROM THE PLOT, THE PRECISION DROPS  
RAPIDLY AS RECALL INCREASES. THIS  
SUGGESTS THAT AS THE MODEL BECOMES  
MORE LENIENT (I.E., INCREASES RECALL),  
IT STARTS MAKING MORE FALSE POSITIVE  
DETECTIONS, LEADING TO LOWER  
PRECISION.

# CONCLUSION

- THE MODEL STRUGGLES TO MAINTAIN PRECISION AS RECALL INCREASES.
- MOST OF THE PRECISION SCORES ARE BELOW 0.4, SUGGESTING A HIGH FALSE POSITIVE RATE.
- LOW F1 SCORES INDICATE A LACK OF BALANCE BETWEEN PRECISION AND RECALL, MEANING THE MODEL MAY EITHER MISS OBJECTS (LOW RECALL) OR MISCLASSIFY NON-OBJECTS AS OBJECTS (LOW PRECISION).
- THE FACT THAT F1 SCORES DON'T RISE SIGNIFICANTLY AT ANY CONFIDENCE LEVEL SUGGESTS THAT OVERALL DETECTION QUALITY IS POOR .

# USING DATA AUGMENTATION : TO IMPROVE THE MODEL

**VERSIONS**

2024-10-10 6:51am  
v2 · a month ago  
 240 640x640 Stretch to

2024-10-04 5:43am  
v1 · 2 months ago  
 100 Fast COCO

**240 Total Images** [View All Images →](#)



**Dataset Split**

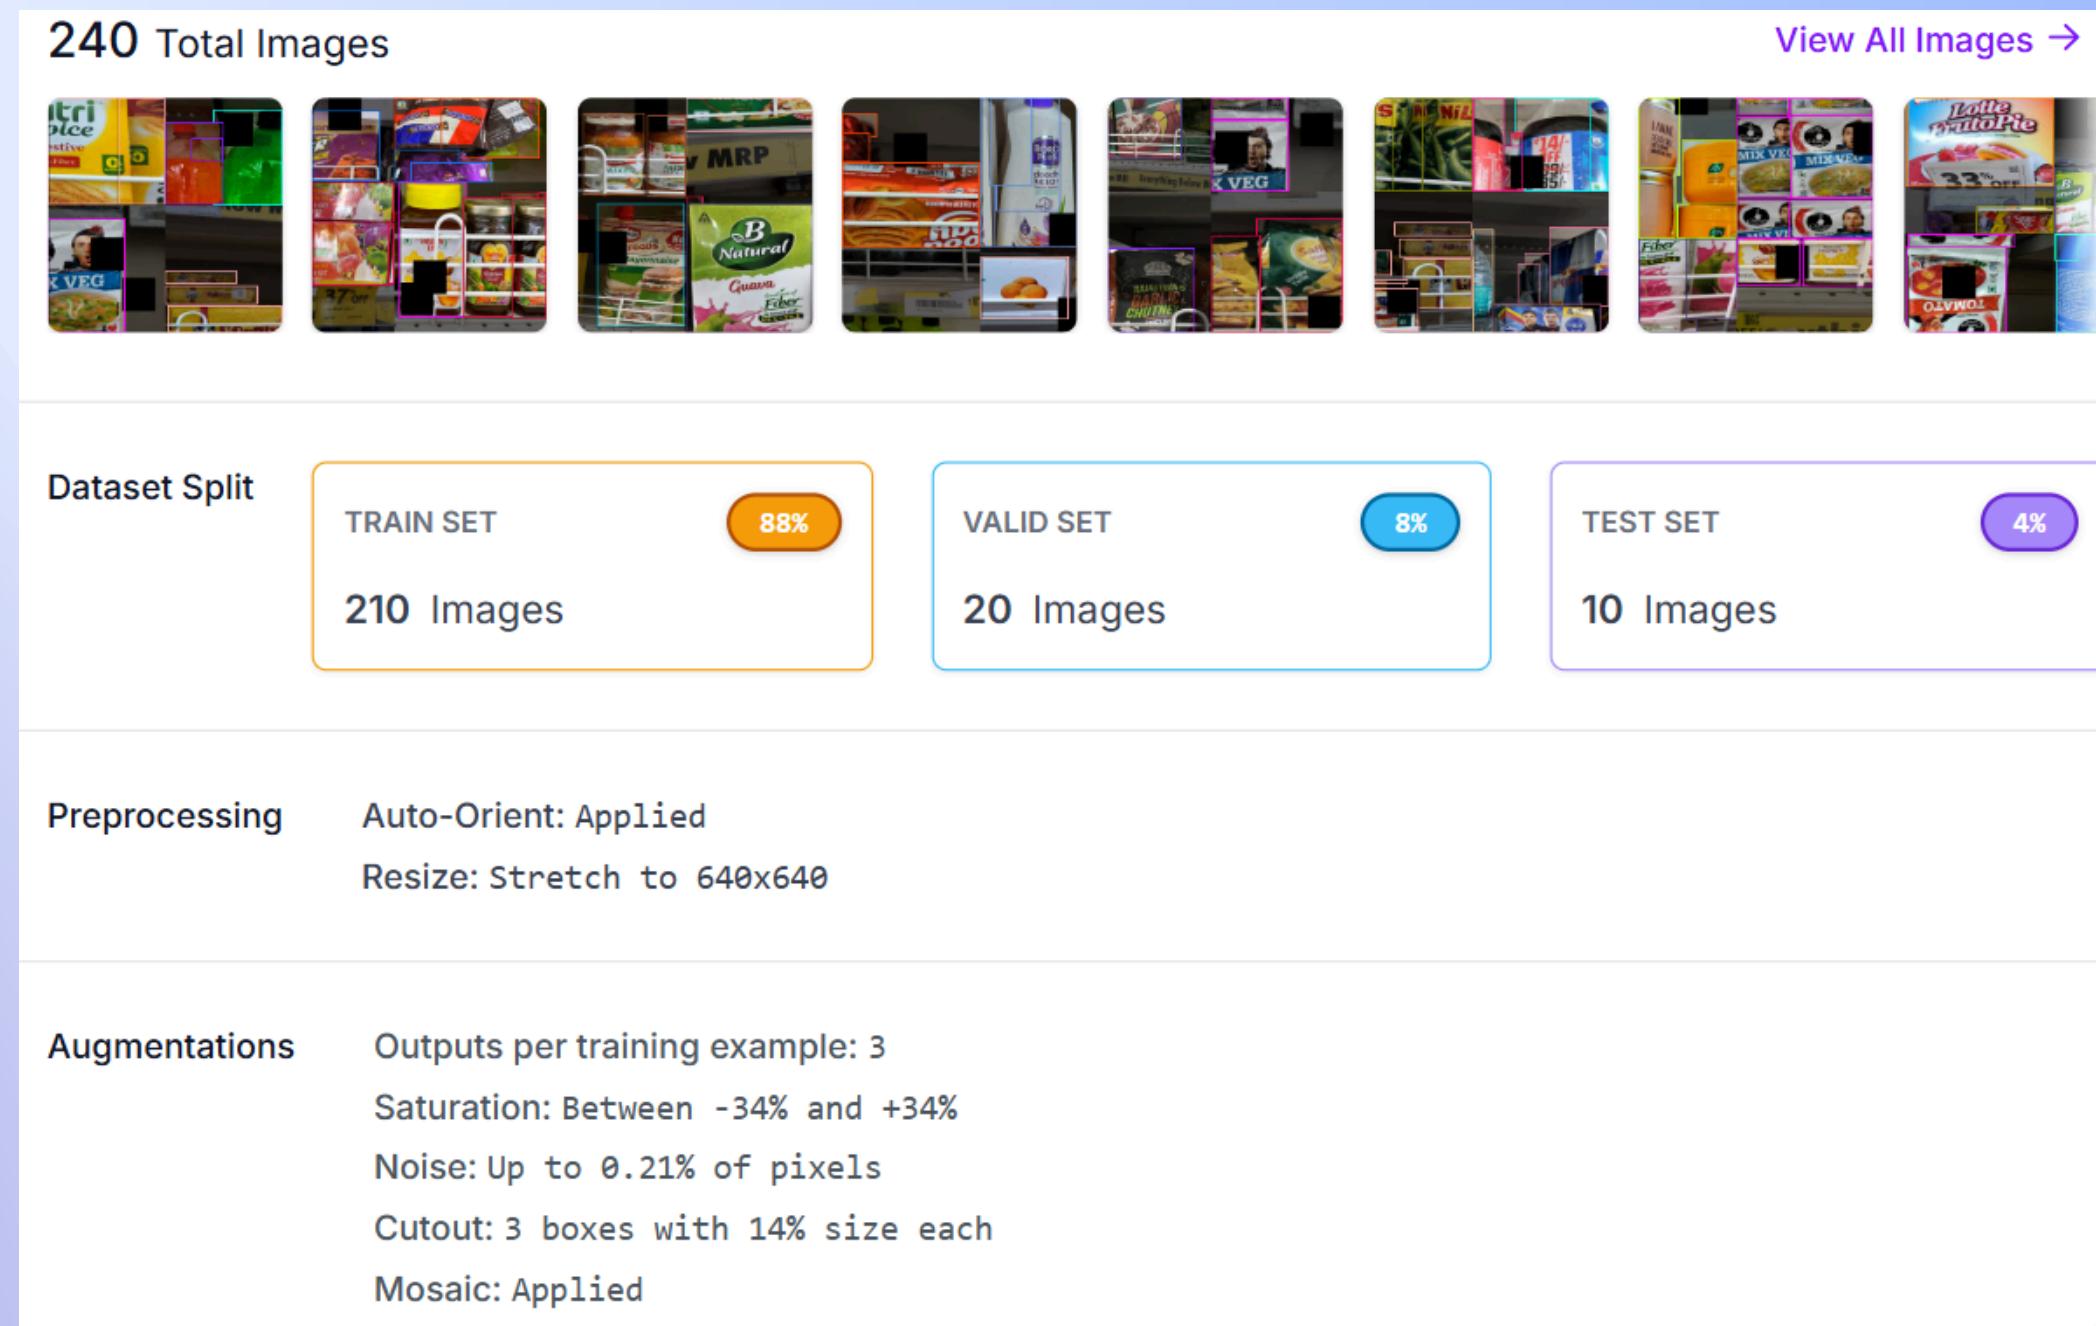
<b>TRAIN SET</b> 210 Images	<b>VALID SET</b> 20 Images	<b>TEST SET</b> 10 Images
--------------------------------	-------------------------------	------------------------------

**Preprocessing** Auto-Orient: Applied  
Resize: Stretch to 640x640

**Augmentations** Outputs per training example: 3  
Saturation: Between -34% and +34%  
Noise: Up to 0.21% of pixels  
Cutout: 3 boxes with 14% size each  
Mosaic: Applied

# CREATING DIFFERENT VERSIONS OF THE MODEL USING DATA AUGMENTATION:

V2:



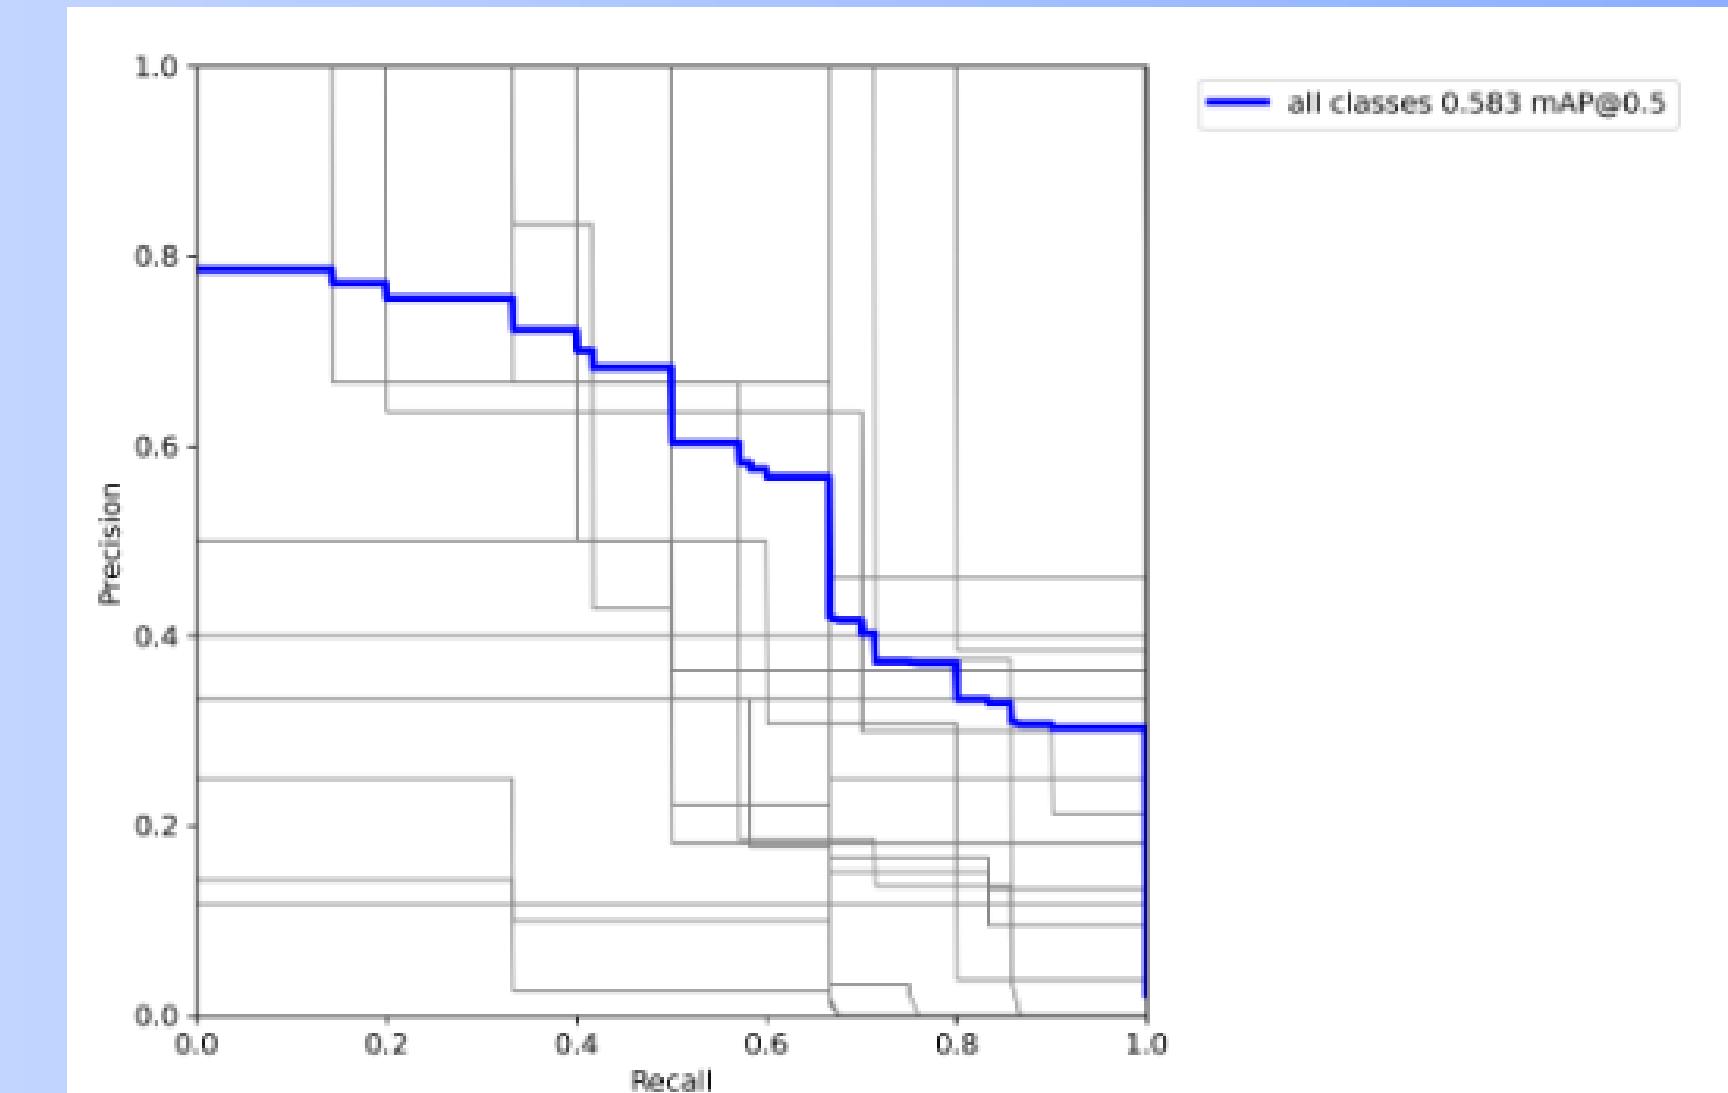
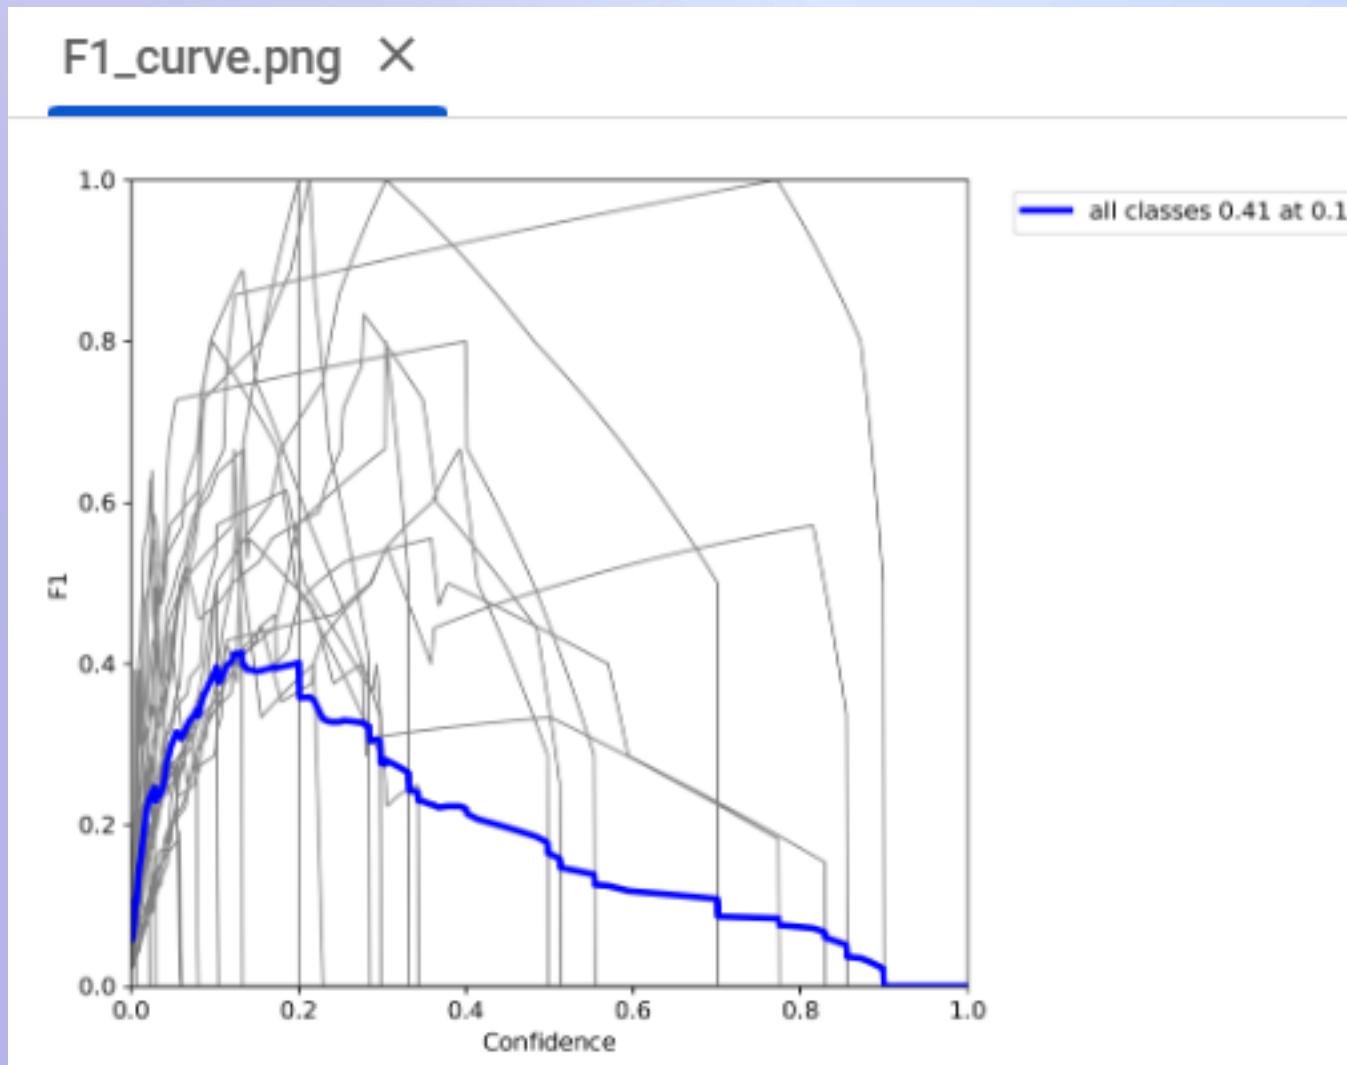
## RESULTS OBTAINED FROM V2:

P  
**0.524**

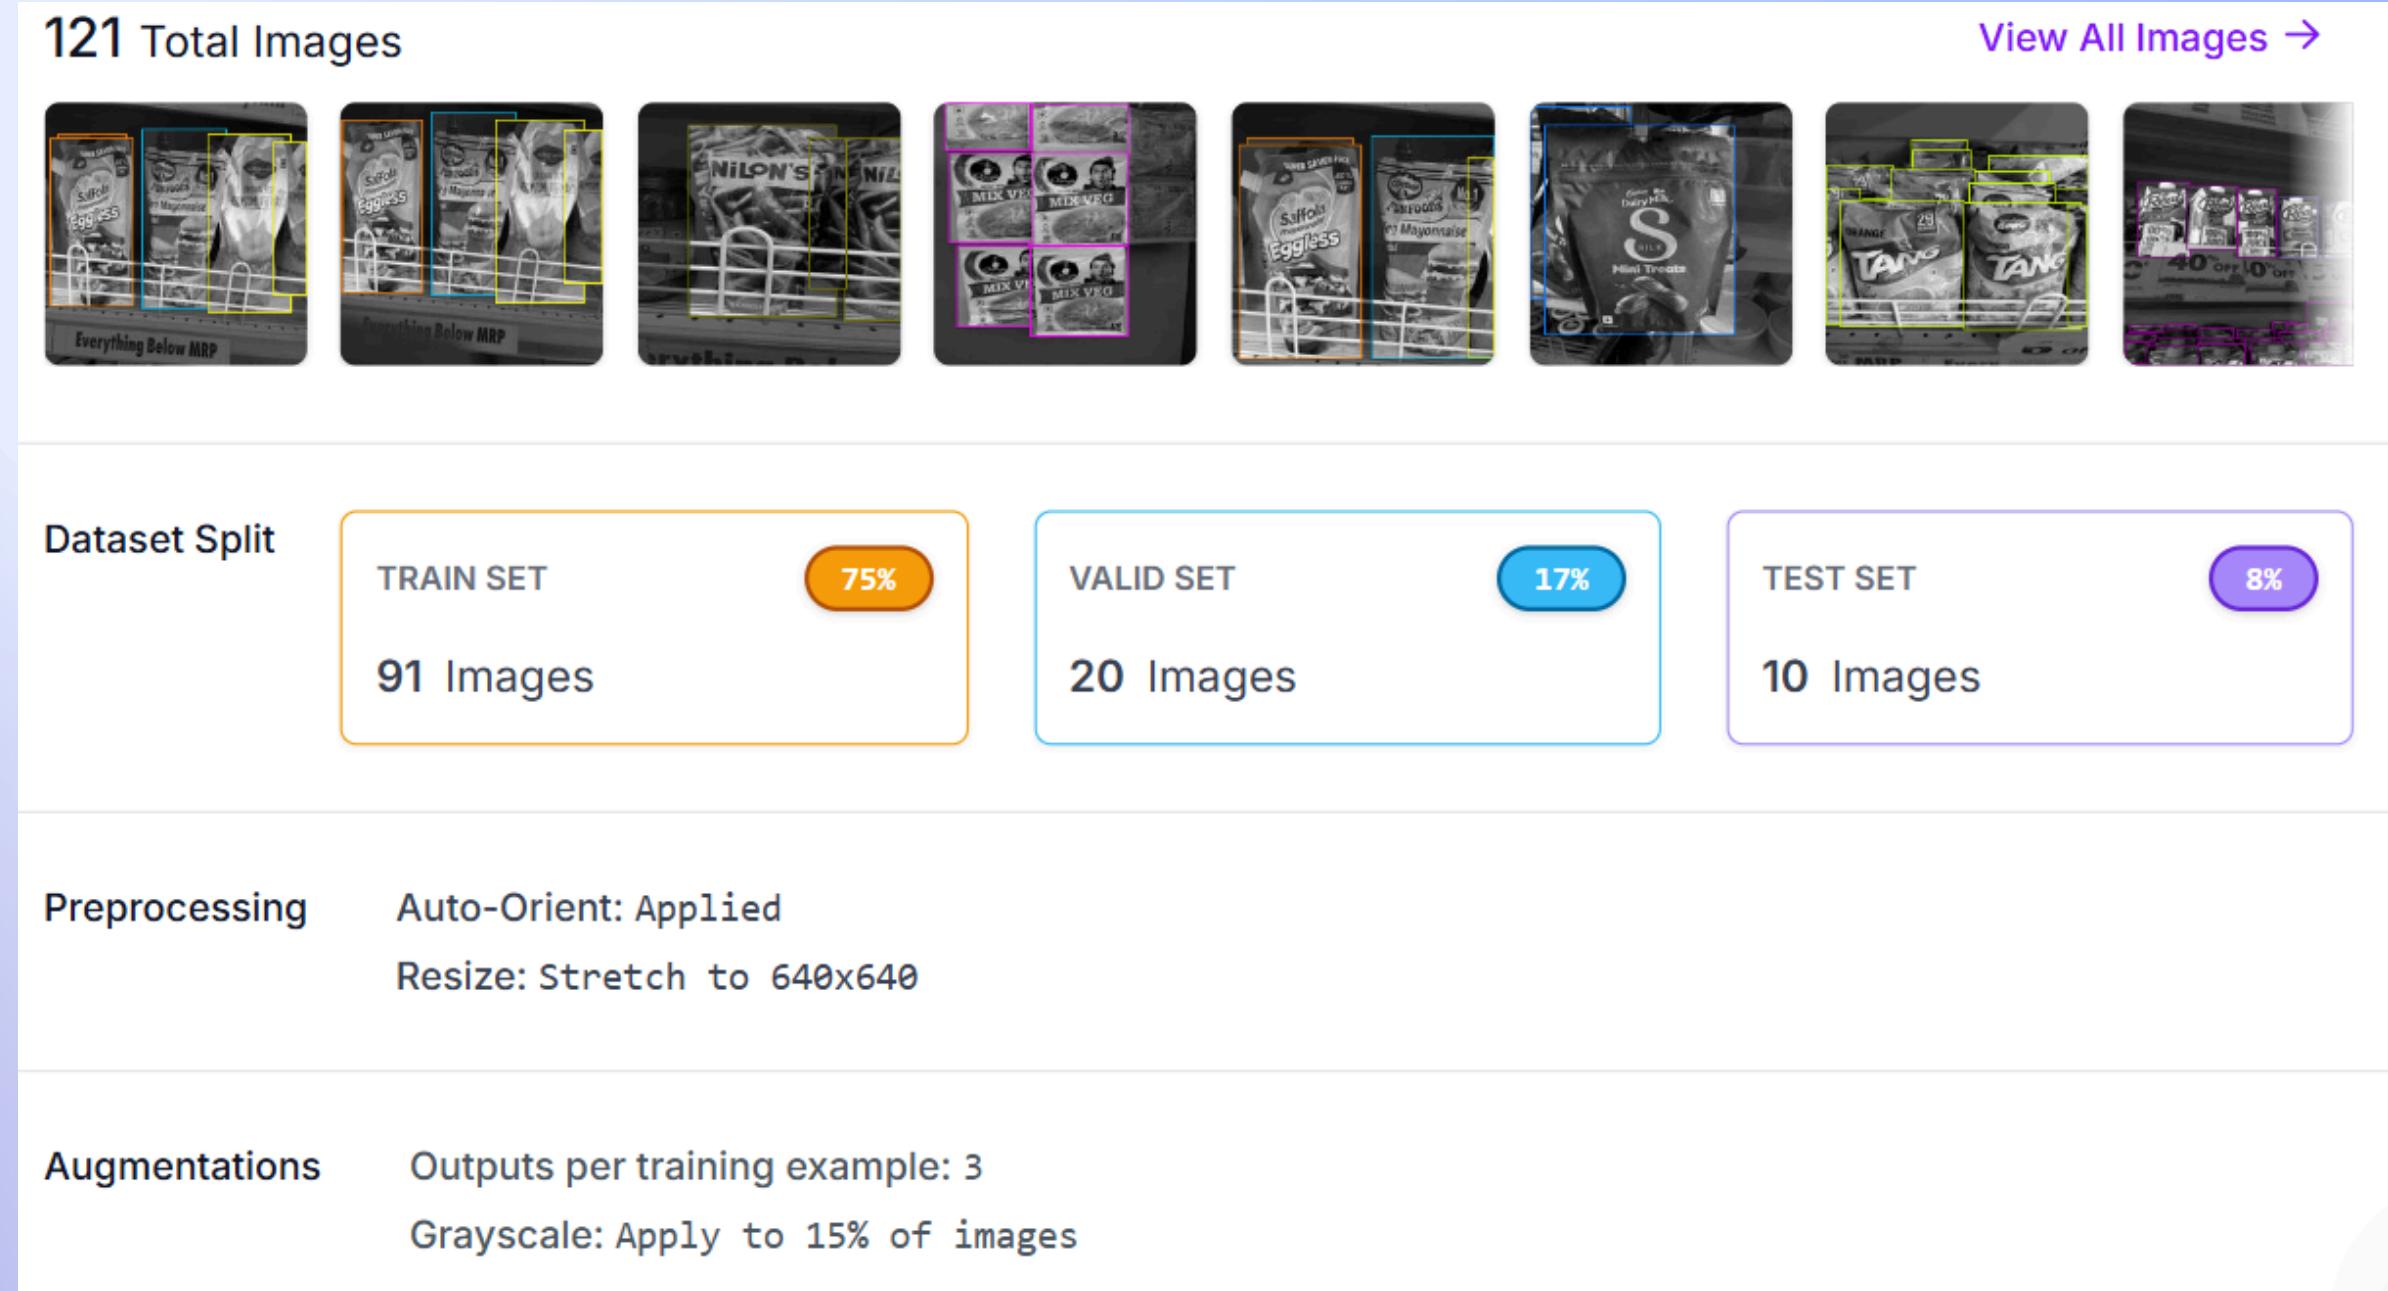
R  
**0.547**

mAP@.5  
**0.578**

mAP@.5 : .95 :  
**0.388**

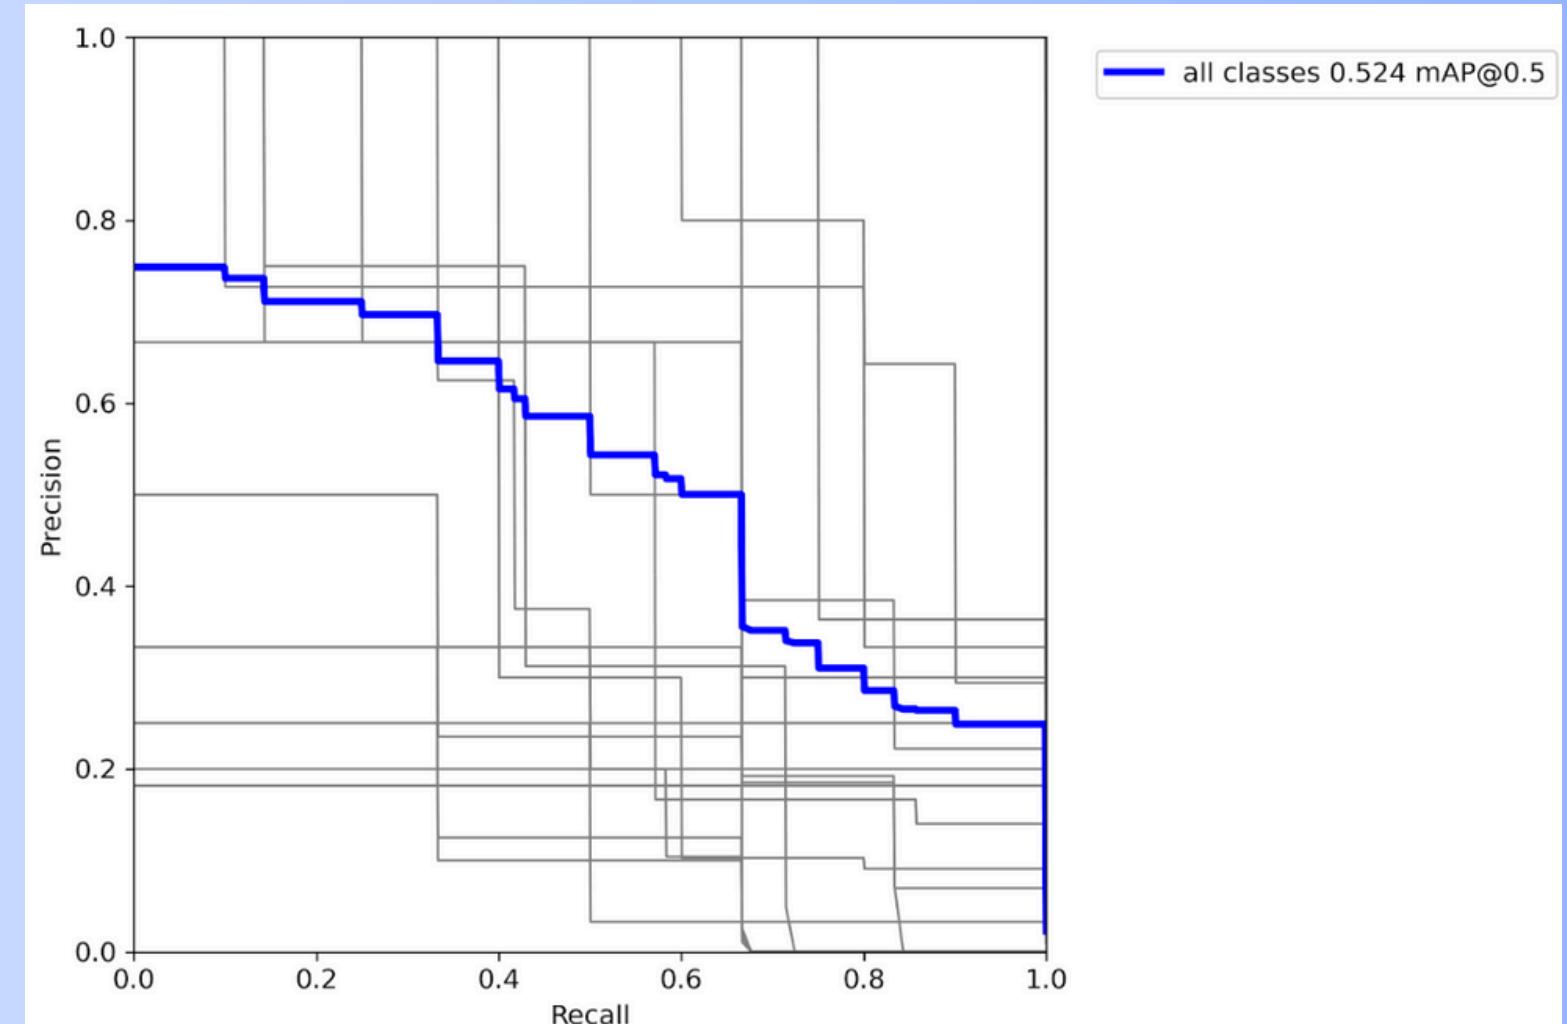
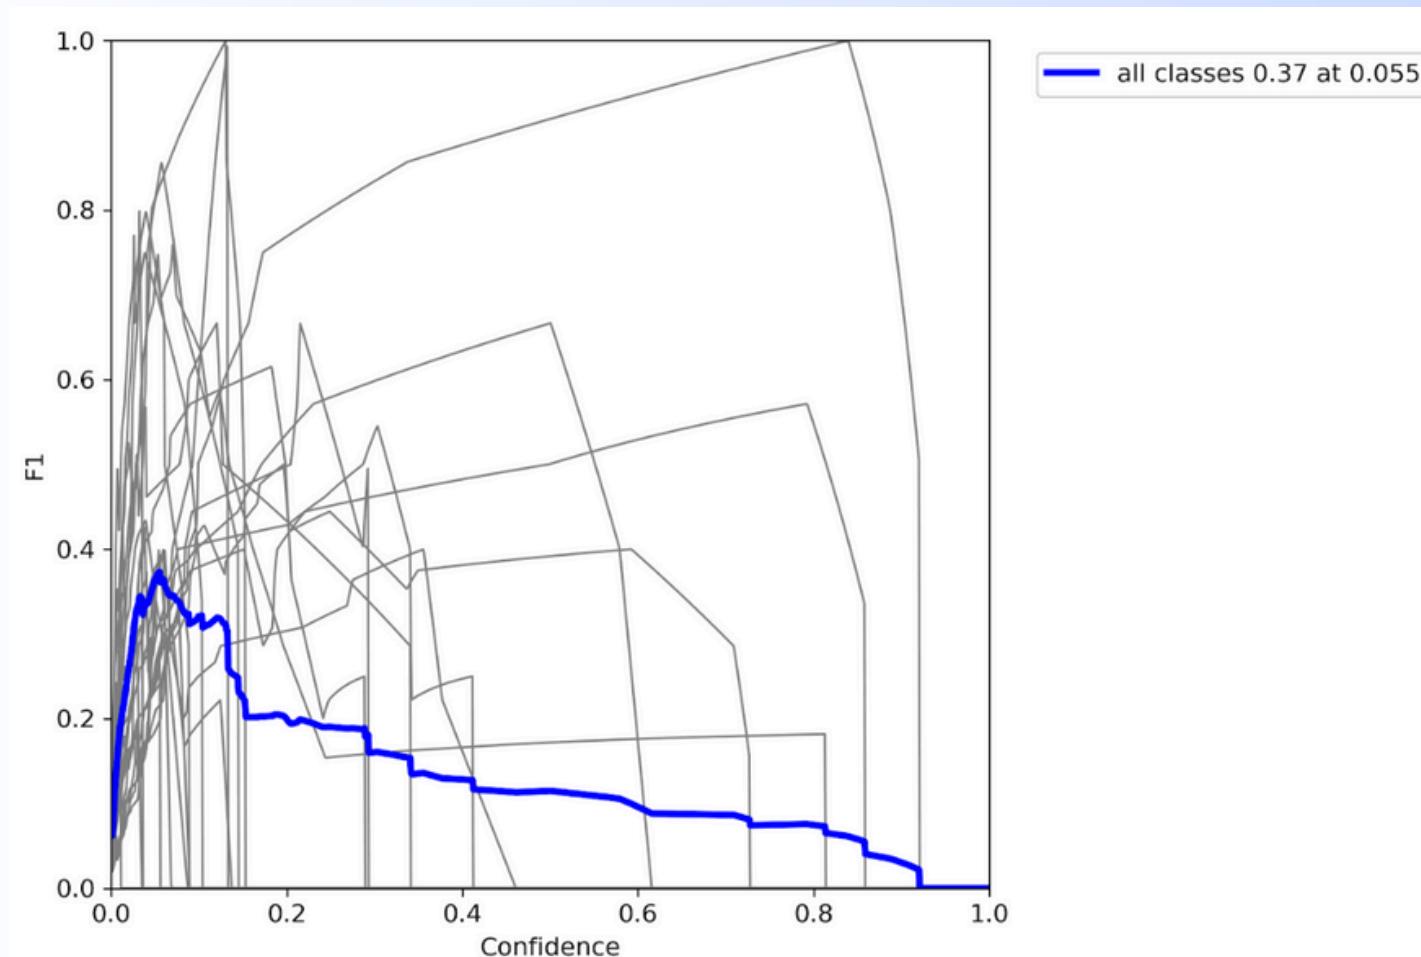


# V3:



# RESULTS OBTAINED FROM V3:

P	R	mAP@.5	mAP@.5 : .95 :
0.771	0.435	0.584	0.396



V4:

240 Total Images

[View All Images →](#)



Dataset Split

TRAIN SET

88%

210 Images

VALID SET

8%

20 Images

TEST SET

4%

10 Images

Preprocessing

Auto-Orient: Applied

Resize: Stretch to 640x640

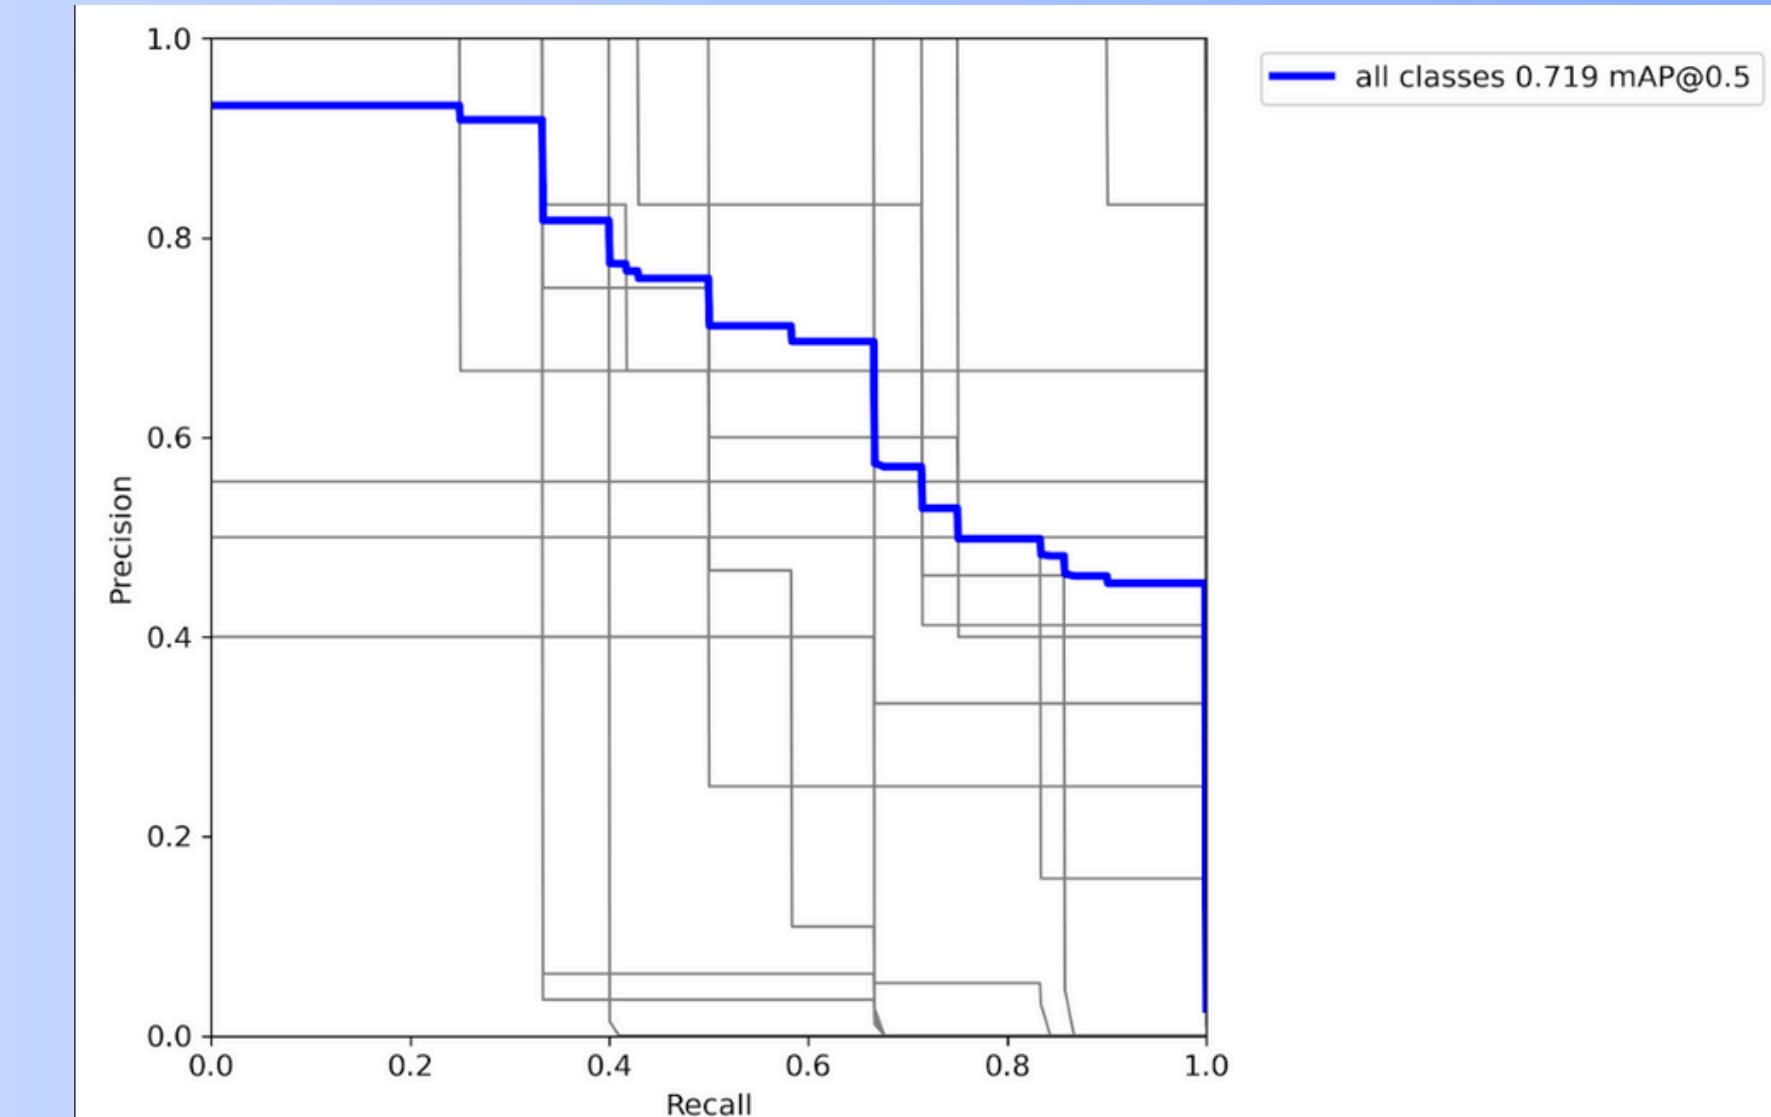
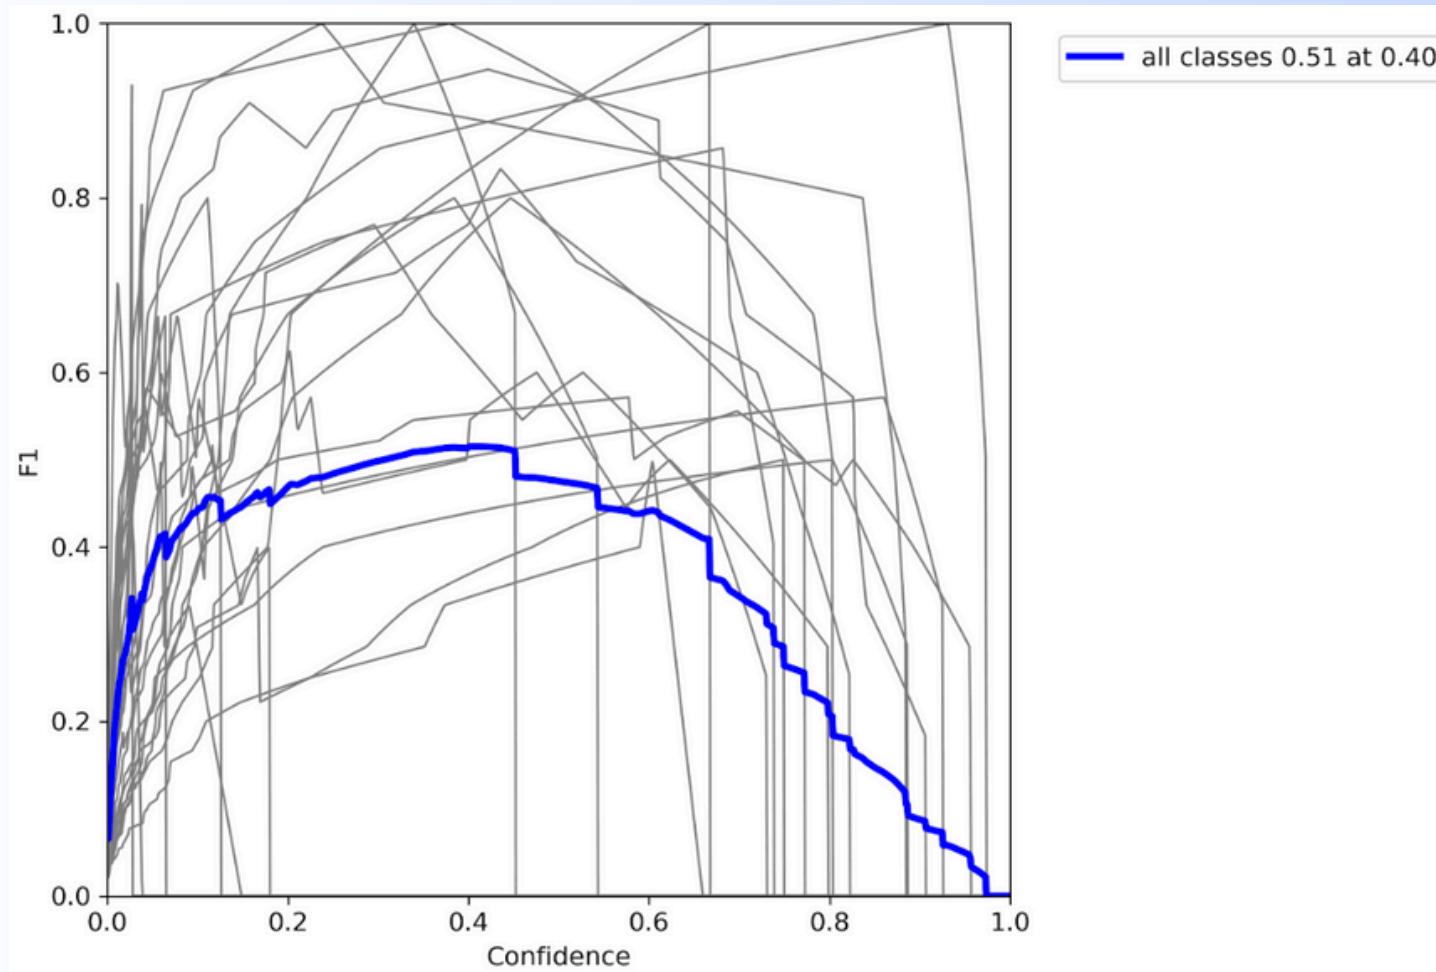
Augmentations

Outputs per training example: 3

Mosaic: Applied

# RESULTS OBTAINED FROM V4:

P	R	mAP@.5	mAP@.5 : .95:
0.792	0.634	0.746	0.531



# CONCLUSION

- SUCCESSFULLY ENHANCED OBJECT DETECTION MODEL THROUGH DATA AUGMENTATION.
- AMONG ALL THE MODELS V4 IS PERFORMING BEST WHICH CAN BE USED FOR OBJECT DETECTION TASK



# **REAL-LIFE USE CASE OF THIS MODEL:**

## **\*RETAIL INVENTORY MANAGEMENT\***

- **\*USE CASE\*: RETAIL STORES CAN USE THE MODEL TO AUTOMATE STOCK COUNTING AND INVENTORY MANAGEMENT. THE MODEL CAN SCAN SHELVES OR STORAGE AREAS TO IDENTIFY AND COUNT THE NUMBER OF PRODUCTS, DETECTING LOW STOCK OR MISPLACED ITEMS.**
- **\*IMPACT\*: REDUCES HUMAN LABOR, MINIMIZES ERRORS, AND ENSURES MORE ACCURATE INVENTORY RECORDS, LEADING TO BETTER STOCK MANAGEMENT, REORDERING, AND AVOIDING STOCKOUTS**

# THANK YOU